

Medical-Evidence-Synthesizer

A Multimodal Diagnostic RAG for Complex Patient Cases

For our final semester long project, we will develop a comprehensive tutorial on building a clinical decision support tool called the "Medical Evidence Synthesizer." This project will guide users through the data science pipeline to create a Retrieval Augmented Generation (RAG) model. The model will ingest complex, de-identified patient data combining structured lab results and unstructured clinical notes to generate a differential diagnosis. A key feature of this tool is its ability to provide verifiable, citable evidence for each diagnostic possibility directly from the source data, showcasing an end-to-end data science solution with a focus on machine learning and explainability.

Structured Data Preprocessing & Cleaning

```
In [ ]: from google.colab import drive  
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call
drive.mount("/content/drive", force_remount=True).

```
In [ ]: %cd /content/drive/MyDrive/MSML_602_FP/Dataset/workspace_data/  
%pwd  
  
/content/drive/MyDrive/MSML_602_FP/Dataset/workspace_data  
Out[ ]: '/content/drive/MyDrive/MSML_602_FP/Dataset/workspace_data'
```

```
In [ ]: !ls  
  
admissions.csv      icustays.csv    omr.csv      processed  
d_icd_diagnoses.csv labevents.csv  patients.csv
```

```
In [ ]: from __future__ import annotations  
  
import json  
from dataclasses import dataclass  
from pathlib import Path  
from typing import Dict, Iterable, List, Tuple  
  
import numpy as np  
import pandas as pd  
  
try:  
    from sklearn.feature_selection import mutual_info_classif  
  
    SKLEARN_AVAILABLE = True  
except ImportError:  
    SKLEARN_AVAILABLE = False
```

```

# Set base directory to the specified folder in Google Drive
BASE_DIR = Path("/content/drive/MyDrive/MSML_602_FP/Dataset/workspace_data")
PROCESSED_DIR = BASE_DIR / "processed"

# Define raw file paths relative to the base directory
RAW_FILES = {
    "admissions": BASE_DIR / "admissions.csv",
    "patients": BASE_DIR / "patients.csv",
    "icustays": BASE_DIR / "icustays.csv",
    "labevents": BASE_DIR / "labevents.csv",
    "omr": BASE_DIR / "omr.csv",
    "diagnoses": BASE_DIR / "d_icd_diagnoses.csv",
}

# Create the processed directory if it doesn't exist
PROCESSED_DIR.mkdir(parents=True, exist_ok=True)

```

In []:

```

from pathlib import Path

# Define input and output paths
base_dir = Path("/content/sample_data")
processed_dir = Path("/content/processed")
processed_dir.mkdir(exist_ok=True)

discharge_input = base_dir / "discharge.csv" # Assuming discharge.csv is in
discharge_output = processed_dir / "discharge_cleaned.csv"

radiology_input = base_dir / "radiology.csv" # Assuming radiology.csv is in
radiology_output = processed_dir / "radiology_cleaned.csv"

# Call the cleaning functions
# Make sure to run the cells defining clean_discharge_notes and clean_radiology_notes
# clean_discharge_notes(discharge_input, discharge_output)
# clean_radiology_notes(radiology_input, radiology_output)

print("Cleaning functions are ready to be called. Uncomment the lines above")

```

Cleaning functions are ready to be called. Uncomment the lines above and run this cell to execute them.

Utility helpers

In []:

```

def _ensure_int(series: pd.Series) -> pd.Series:
    """Return a nullable integer series; preserves missing values."""

    return pd.to_numeric(series, errors="coerce").astype("Int64")


def _safe_to_numeric(series: pd.Series) -> pd.Series:
    """Convert object series to float when possible, otherwise NaN."""

    if series.dtype.kind in {"i", "u", "f"}:
        return series
    return pd.to_numeric(series.astype(str).str.replace(",",""), errors="coerce")

```

```

def _duration_hours(end: pd.Series, start: pd.Series) -> pd.Series:
    """Compute duration in hours between two datetime series."""
    return (end - start).dt.total_seconds() / 3600.0

def _duration_minutes(end: pd.Series, start: pd.Series) -> pd.Series:
    return (end - start).dt.total_seconds() / 60.0

def _recent_records(df: pd.DataFrame, group_cols: List[str], timestamp_col: str):
    """Return the most recent row per group based on timestamp_col."""
    idx = df.groupby(group_cols)[timestamp_col].transform("idxmax")
    return df.loc[idx].reset_index(drop=True)

def _save_table(df: pd.DataFrame, filename: str) -> str:
    """Persist DataFrame as parquet if available; otherwise CSV.

    Returns the final filename that was written.
    """
    target = PROCESSED_DIR / filename
    suffix = target.suffix.lower()
    if suffix == ".parquet":
        try:
            df.to_parquet(target, index=False)
            return target.name
        except ImportError:
            csv_target = target.with_suffix(".csv")
            df.to_csv(csv_target, index=False)
            return csv_target.name
    elif suffix == ".csv":
        df.to_csv(target, index=False)
        return target.name
    else:
        raise ValueError(f"Unsupported file extension for {target}")

```

Admissions

```
In [ ]: def load_and_clean_admissions(path: Path) -> pd.DataFrame:
    df = pd.read_csv(
        path,
        parse_dates=[
            "admittime",
            "dischtime",
            "deathtime",
            "edregtime",
            "edouttime",
        ],
        na_values=["", "NA", "NaN", "?"],
    )

    df = df.drop_duplicates(subset=["hadm_id"], keep="last")
```

```

df["subject_id"] = _ensure_int(df["subject_id"])
df["hadm_id"] = _ensure_int(df["hadm_id"])

# Feature engineering
df["los_hours"] = _duration_hours(df["dischtime"], df["admittime"])
df.loc[df["los_hours"] < 0, "los_hours"] = np.nan # guard bad timestamp

df["ed_wait_minutes"] = _duration_minutes(df["admittime"], df["edregtime"])
df["ed_stay_minutes"] = _duration_minutes(df["edouttime"], df["edregtime"])

df["died_in_hospital"] = df["hospital_expire_flag"].fillna(0).astype("Int64")

# Harmonise key categoricals
categorical_cols = [
    "admission_type",
    "admission_location",
    "discharge_location",
    "insurance",
    "language",
    "marital_status",
    "race",
]
for col in categorical_cols:
    df[col] = (
        df[col]
        .astype(str)
        .str.strip()
        .str.upper()
        .replace({"NAN": np.nan, "?": np.nan, "": np.nan})
    )

df["marital_status"] = df["marital_status"].fillna("UNKNOWN")
df["language"] = df["language"].fillna("UNKNOWN")

return df

```

Patients

```

In [ ]: def load_and_clean_patients(path: Path) -> pd.DataFrame:
    df = pd.read_csv(path, parse_dates=["dod"], na_values=["", "NA", "?"])
    df = df.drop_duplicates(subset=["subject_id"], keep="last")

    df["subject_id"] = _ensure_int(df["subject_id"])
    df["anchor_age"] = pd.to_numeric(df["anchor_age"], errors="coerce")
    df["anchor_year"] = pd.to_numeric(df["anchor_year"], errors="coerce")

    df["dod_available"] = df["dod"].notna().astype("Int64")

    age_bins = [0, 40, 60, 75, 200]
    age_labels = ["<40", "40-59", "60-74", "75+"]
    df["age_group"] = pd.cut(df["anchor_age"], bins=age_bins, labels=age_labels)

    df["gender"] = df["gender"].str.upper().str.strip()

return df

```

ICU stays

```
In [ ]: def load_and_aggregate_icustays(path: Path) -> pd.DataFrame:
    df = pd.read_csv(
        path,
        parse_dates=["intime", "outtime"],
        na_values=["", "NA", "NaN"],
    )

    df["subject_id"] = _ensure_int(df["subject_id"])
    df["hadm_id"] = _ensure_int(df["hadm_id"])
    df["stay_id"] = _ensure_int(df["stay_id"])

    df["los"] = pd.to_numeric(df["los"], errors="coerce")

    aggregations = {
        "stay_id": "count",
        "los": ["sum", "mean", "max"],
    }
    grouped = df.groupby("hadm_id").agg(aggregations)
    grouped.columns = [
        "icu_stay_count",
        "icu_los_hours_sum",
        "icu_los_hours_mean",
        "icu_los_hours_max",
    ]
    first_last_units = df.sort_values("intime").groupby("hadm_id").agg(
        first_careunit_first=("first_careunit", "first"),
        first_careunit_last=("first_careunit", "last"),
        last_careunit_first=("last_careunit", "first"),
        last_careunit_last=("last_careunit", "last"),
    )
    aggregated = grouped.join(first_last_units, how="left")
    aggregated.reset_index(inplace=True)

    return aggregated
```

Lab events

```
In [ ]: @dataclass
class LabFeatureSpec:
    itemid: int
    name: str

    def _select_lab_items(
        df: pd.DataFrame, eligible_hadm: Iterable[int], top_n: int = 15, min_cov
    ) -> List[LabFeatureSpec]:
        hadm_set = pd.Index(eligible_hadm)
        coverage = (
            df.dropna(subset=["hadm_id"])
                .groupby("itemid")["hadm_id"]
```

```

        .nunique()
        .sort_values(ascending=False)
    )
coverage_ratio = coverage / hadm_set.nunique()

selected_itemids = coverage_ratio[coverage_ratio >= min_coverage].head(t
specs = [LabFeatureSpec(int(item), f"lab_{item}") for item in selected_i
return specs

def load_and_aggregate_labs(path: Path, eligible_hadm: Iterable[int]) -> Tup
df = pd.read_csv(
    path,
    parse_dates=["charttime", "storetime"],
    na_values=["", "NA", "NaN", "?", "__"],
)

df["subject_id"] = _ensure_int(df["subject_id"])
df["hadm_id"] = _ensure_int(df["hadm_id"])

# Prefer valuenum; fall back to parsed value
df["value_num_clean"] = df["valuenum"]
missing_mask = df["value_num_clean"].isna()
df.loc[missing_mask, "value_num_clean"] = _safe_to_numeric(df.loc[missin

df = df.dropna(subset=["hadm_id", "value_num_clean", "charttime"])

specs = _select_lab_items(df, eligible_hadm)
if not specs:
    return pd.DataFrame(columns=["hadm_id"], [])

frames = []
for spec in specs:
    subset = df[df["itemid"] == spec.itemid].copy()
    if subset.empty:
        continue
    subset.sort_values("charttime", inplace=True)

    agg = subset.groupby("hadm_id").agg(
        **{
            f"{spec.name}_count": ("value_num_clean", "count"),
            f"{spec.name}_mean": ("value_num_clean", "mean"),
            f"{spec.name}_std": ("value_num_clean", "std"),
            f"{spec.name}_min": ("value_num_clean", "min"),
            f"{spec.name}_max": ("value_num_clean", "max"),
        }
    )

    last_values = (
        subset.sort_values("charttime")
        .groupby("hadm_id", as_index=False)
        .tail(1)
        .set_index("hadm_id")["value_num_clean"]
    )
    agg[f"{spec.name}_last"] = agg.index.map(last_values)

```

```

        frames.append(agg)

    if not frames:
        return pd.DataFrame(columns=["hadm_id"]), []

    lab_features = pd.concat(frames, axis=1)
    lab_features.reset_index(inplace=True)
    lab_features = lab_features.loc[:, ~lab_features.columns.duplicated()]

    return lab_features, specs

```

Outpatient (OMR) measurements

```

In [ ]: def _parse_blood_pressure(value: str) -> Tuple[float | None, float | None]:
    if not isinstance(value, str):
        return (np.nan, np.nan)
    parts = value.replace(" ", "").split("/")
    if len(parts) != 2:
        return (np.nan, np.nan)
    systolic = pd.to_numeric(parts[0], errors="coerce")
    diastolic = pd.to_numeric(parts[1], errors="coerce")
    return (systolic, diastolic)

def load_and_aggregate_omr(path: Path) -> pd.DataFrame:
    df = pd.read_csv(path, parse_dates=["chartdate"], na_values=["", "NA", ""])
    df["subject_id"] = _ensure_int(df["subject_id"])

    df["result_name"] = df["result_name"].str.strip().str.upper()

    numeric_mask = df["result_name"].isin(
        [
            "HEIGHT (INCHES)",
            "WEIGHT (LBS)",
            "BMI (KG/M2)",
            "WEIGHT (LBS)",
        ]
    )
    df.loc[numeric_mask, "result_value_numeric"] = _safe_to_numeric(df.loc[r
latest = _recent_records(df, ["subject_id", "result_name"], "chartdate")

# Convert to feature columns
features: Dict[str, pd.Series] = {"subject_id": latest["subject_id"]}

def _assign_feature(name: str, mask: pd.Series, values: pd.Series) -> N
    colname = f"omr_{name}"
    features[colname] = values.where(mask).groupby(latest["subject_id"])

height_mask = latest["result_name"] == "HEIGHT (INCHES)"
height_cm = latest.loc[height_mask, "result_value_numeric"] * 2.54
_assign_feature("height_cm", height_mask, height_cm)

weight_mask = latest["result_name"] == "WEIGHT (LBS)"
weight_kg = latest.loc[weight_mask, "result_value_numeric"] * 0.45359237

```

```

    _assign_feature("weight_kg", weight_mask, weight_kg)

    bmi_mask = latest["result_name"] == "BMI (KG/M2)"
    bmi_val = latest.loc[bmi_mask, "result_value_numeric"]
    _assign_feature("bmi", bmi_mask, bmi_val)

    bp_mask = latest["result_name"] == "BLOOD PRESSURE"
    bp_vals = latest.loc[bp_mask, "result_value"].apply(_parse_blood_pressure)
    if not bp_vals.empty:
        systolic = bp_vals.apply(lambda x: x[0])
        diastolic = bp_vals.apply(lambda x: x[1])
        _assign_feature("blood_pressure_systolic", bp_mask, systolic)
        _assign_feature("blood_pressure_diastolic", bp_mask, diastolic)

    result = pd.DataFrame(features).groupby("subject_id").first().reset_index()
    return result

```

Feature assembly & selection

```

In [ ]: import pandas as pd
import numpy as np
from typing import List, Tuple, Dict
from dataclasses import dataclass

# Assuming mutual_info_classif is needed and SKLEARN_AVAILABLE is defined elsewhere
# from sklearn.feature_selection import mutual_info_classif

def assemble_feature_table(
    admissions: pd.DataFrame,
    patients: pd.DataFrame,
    icu: pd.DataFrame,
    lab: pd.DataFrame,
    omr: pd.DataFrame,
) -> pd.DataFrame:
    df = admissions.merge(patients, on="subject_id", how="left", suffixes=("",))

    if not icu.empty:
        df = df.merge(icu, on="hadm_id", how="left")

    if not lab.empty:
        df = df.merge(lab, on="hadm_id", how="left")

    if not omr.empty:
        df = df.merge(omr, on="subject_id", how="left")

    # Drop columns with excessive missingness (> 70%)
    missing_ratio = df.isna().mean()
    keep_cols = missing_ratio[missing_ratio <= 0.7].index.tolist()
    df = df[keep_cols]

    return df

def _prepare_numeric_matrix(df: pd.DataFrame, target_col: str) -> Tuple[pd.D
feature_df = df.drop(columns=[target_col]).copy()

```

```

target = df[target_col].astype("Int64").fillna(0)

# Exclude identifiers from modeling features
identifier_cols = [col for col in feature_df.columns if col in {"subject_id", "hadm_id"}]
feature_df = feature_df.drop(columns=identifier_cols, errors="ignore")

# Convert datetime and timedelta columns to numeric representations
datetime_cols = feature_df.select_dtypes(include=["datetime64[ns]", "date"])
for col in datetime_cols:
    feature_df[col] = feature_df[col].apply(lambda x: x.value if pd.notnull(x) else 0)

timedelta_cols = feature_df.select_dtypes(include=["timedelta64[ns]"])
for col in timedelta_cols:
    feature_df[col] = feature_df[col].dt.total_seconds()

# One-hot encode categorical variables using a potentially more memory-efficient approach
categorical_cols = feature_df.select_dtypes(include=["object", "category"])
# Use get_dummies with sparse=True if available in your pandas version as it is faster
# feature_df = pd.get_dummies(feature_df, columns=categorical_cols, dummy_na=True)
# For broader compatibility, stick to standard get_dummies but be mindful of memory usage
feature_df = pd.get_dummies(feature_df, columns=categorical_cols, dummy_na=True)

# Impute missing values *after* one-hot encoding to handle new dummy columns
# Use median imputation for numeric columns
for col in feature_df.select_dtypes(include=[np.number]).columns:
    if feature_df[col].isnull().any():
        feature_df[col].fillna(feature_df[col].median(), inplace=True)

# Drop constant columns
nunique = feature_df.nunique()
feature_df = feature_df.loc[:, nunique > 1]

return feature_df, target

def select_features(df: pd.DataFrame, target_col: str, k: int = 20) -> Tuple[pd.DataFrame, np.ndarray]:
    global SKLEARN_AVAILABLE # Declare SKLEARN_AVAILABLE as global
    X, y = _prepare_numeric_matrix(df, target_col)

    if X.empty:
        return df[["subject_id", "hadm_id", target_col]].copy(), [], {}

    feature_scores: Dict[str, float] = {}

    # Check if SKLEARN_AVAILABLE and mutual_info_classif are available before proceeding
    # Assuming SKLEARN_AVAILABLE is a boolean defined in a previous cell
    # Assuming mutual_info_classif is imported in a previous cell if SKLEARN_AVAILABLE is True
    if 'SKLEARN_AVAILABLE' in globals() and SKLEARN_AVAILABLE and y.nunique() > 2:
        try:
            from sklearn.feature_selection import mutual_info_classif
            scores = mutual_info_classif(X, y)
            feature_scores = dict(zip(X.columns, scores))
        except ImportError:
            print("Scikit-learn not available or mutual_info_classif not imported")
            SKLEARN_AVAILABLE = False # Update flag if import fails

```

```

corrs = {}
for col in X.columns:
    series = X[col]
    values = series.to_numpy(dtype=float)
    if np.nanstd(values) == 0:
        corrs[col] = 0.0
        continue
    # Ensure y has no NaNs for correlation calculation
    valid_indices = ~np.isnan(values)
    if np.sum(valid_indices) > 1: # Need at least two non-NaN values
        corrs[col] = abs(np.corrcoef(values[valid_indices], y[valid_indices])[0, 1])
    else:
        corrs[col] = 0.0 # Cannot compute correlation with less than 2 values

feature_scores = corrs

else:
    # Fallback: absolute Pearson correlation
    corrs = {}
    for col in X.columns:
        series = X[col]
        values = series.to_numpy(dtype=float)
        if np.nanstd(values) == 0:
            corrs[col] = 0.0
            continue
        # Ensure y has no NaNs for correlation calculation
        valid_indices = ~np.isnan(values)
        if np.sum(valid_indices) > 1: # Need at least two non-NaN values
            corrs[col] = abs(np.corrcoef(values[valid_indices], y[valid_indices])[0, 1])
        else:
            corrs[col] = 0.0 # Cannot compute correlation with less than 2 values

feature_scores = corrs

# Ensure feature_scores is not empty before sorting
if not feature_scores:
    return df[["subject_id", "hadm_id", target_col]].copy(), [], {}

top_features = sorted(feature_scores, key=feature_scores.get, reverse=True)

# Ensure selected_df is created correctly even if top_features is empty
if top_features:
    selected_df = df[["subject_id", "hadm_id", target_col]].join(X[top_features])
else:
    selected_df = df[["subject_id", "hadm_id", target_col]].copy()

return selected_df, top_features, feature_scores

```

Main orchestration

```
In [ ]: from pathlib import Path

def main() -> None:
```

```

# Use the BASE_DIR defined in cell 9cG3wUVrmBFQ
# BASE_DIR = Path(".").resolve()
PROCESSED_DIR = BASE_DIR / "processed"
PROCESSED_DIR.mkdir(exist_ok=True)

# Check if features_selected.csv already exists
features_selected_path = PROCESSED_DIR / "features_cleaned_selected.csv"
if features_selected_path.exists():
    print(f'{features_selected_path} found. Skipping data processing')
    # Optionally load the existing features_selected.csv here if needed
    # features_selected_df = pd.read_csv(features_selected_path)
    return # Exit the function if the file exists

admissions = load_and_clean_admissions(BASE_DIR / RAW_FILES["admissions"])
patients = load_and_clean_patients(BASE_DIR / RAW_FILES["patients"])
icu = load_and_aggregate_icustays(BASE_DIR / RAW_FILES["icustays"])

# Commenting out processing of lab events and feature selection to prevent
lab_path = BASE_DIR / RAW_FILES["labevents"]
labs = pd.DataFrame()
lab_specs: List[LabFeatureSpec] = []
if lab_path.exists():
    print(f'Processing lab events from {lab_path} in chunks...')
    # Process labevents in chunks
    chunk_size = 100000 # Adjust chunk size based on available RAM and time
    eligible_hadm = admissions["hadm_id"].dropna().unique()
    all_lab_features = []
    first_chunk = True
    for chunk in pd.read_csv(lab_path, chunksize=chunk_size, on_bad_lines='skip',
                             parse_dates=["charttime", "storetime"], na_values="NA"):
        chunk["subject_id"] = _ensure_int(chunk["subject_id"])
        chunk["hadm_id"] = _ensure_int(chunk["hadm_id"])

        # Prefer valuenum; fall back to parsed value
        chunk["value_num_clean"] = chunk["valuenum"]
        missing_mask = chunk["value_num_clean"].isna()
        chunk.loc[missing_mask, "value_num_clean"] = _safe_to_numeric(chunk,
                                                                     "value_num_clean",
                                                                     "value_num_clean")

        chunk = chunk.dropna(subset=["hadm_id", "value_num_clean", "charttime"])

    if first_chunk:
        # Select lab items based on the first chunk (or a representative subset)
        # For better accuracy, might need a separate pass or larger chunk
        lab_specs = _select_lab_items(chunk, eligible_hadm, top_n=50)
        first_chunk = False

    if not lab_specs:
        print("No eligible lab items found in the initial chunk. Skipped")
        break # Exit loop if no specs are found

    frames = []
    for spec in lab_specs:
        subset = chunk[chunk["itemid"] == spec.itemid].copy()
        if subset.empty:

```

```

        continue
subset.sort_values("charttime", inplace=True)

agg = subset.groupby("hadm_id").agg(
    **{
        f"{spec.name}_count": ("value_num_clean", "count"),
        f"{spec.name}_mean": ("value_num_clean", "mean"),
        f"{spec.name}_std": ("value_num_clean", "std"),
        f"{spec.name}_min": ("value_num_clean", "min"),
        f"{spec.name}_max": ("value_num_clean", "max"),
    }
)

last_values = (
    subset.sort_values("charttime")
    .groupby("hadm_id", as_index=False)
    .tail(1)
    .set_index("hadm_id")["value_num_clean"]
)
agg[f"{spec.name}_last"] = agg.index.map(last_values)

frames.append(agg)

if frames:
    chunk_lab_features = pd.concat(frames, axis=1)
    chunk_lab_features.reset_index(inplace=True)
    chunk_lab_features = chunk_lab_features.loc[:, ~chunk_lab_f
    all_lab_features.append(chunk_lab_features)
print(f"Processed a lab chunk.")

if all_lab_features:

    labs = pd.concat(all_lab_features, ignore_index=True)

    labs = labs.groupby('hadm_id', as_index=False).agg('first') # Si

omr_path = BASE_DIR / RAW_FILES["omr"]
omr = pd.DataFrame() # Initialize omr as empty DataFrame
if omr_path.exists():
    print(f"Processing OMR data from {omr_path}...")
    # Uncommenting OMR processing
    omr = load_and_aggregate_omr(omr_path)
    print("OMR processing complete.")
else:
    print(f"OMR file not found at {omr_path}. Skipping OMR processing.")

# Persist intermediate datasets (only admissions, patients, icu, and now
saved_files = {
    "admissions": _save_table(admissions, "clean_admissions.parquet"),
    "patients": _save_table(patients, "clean_patients.parquet"),
    "icu": _save_table(icu, "icu_aggregates.parquet"),
    # "labs": _save_table(labs, "lab_aggregates.parquet") if not labs.empty
    "omr": _save_table(omr, "omr_latest.parquet") if not omr.empty else
}

```

```

# Commenting out feature assembly and selection
feature_table = assemble_feature_table(admissions, patients, icu, labs,
                                       skip_assembly=True)

# Handle case where feature_table might be empty after filtering
if not feature_table.empty:
    selected_df, top_features, feature_scores = select_features(feature_table)
else:
    selected_df = pd.DataFrame()
    top_features = []
    feature_scores = {}
    print("Feature table is empty after assembly. Skipping feature selection")

full_features_filename = _save_table(feature_table, "features_full.csv")
selected_features_filename = _save_table(selected_df, "features_selected.csv")

# Creating dummy metadata to allow the cell to complete
metadata = {
    "lab_features": [], # Empty as lab processing is skipped
    "selected_feature_names": [], # Empty as feature selection is skipped
    "feature_scores": {}, # Empty as feature selection is skipped
    "sklearn_used": False, # False as sklearn is not used in the skipped steps
    "rows": [
        {"feature_full": 0, # 0 as feature assembly is skipped
         "feature_selected": 0, # 0 as feature selection is skipped}
    ],
    "saved_files": saved_files # Includes admissions, patients, icu, and feature tables
}

with open(PREPROCESSED_DIR / "feature_metadata.json", "w", encoding="utf-8") as f:
    json.dump(metadata, f, indent=2)

print("Data preprocessing (partial) complete. Outputs saved under 'processed' directory")
print("Note: Processing of lab events, feature assembly, and feature selection has been skipped.")


main()

```

'/content/drive/MyDrive/MSML_602_FP/Dataset/workspace_data/processed/features_cleaned_selected.csv' found. Skipping data processing and feature selection.

Unstructured Data Preprocessing & Cleaning

Discharge Notes

```
In [ ]: import re
import pandas as pd

def clean_text(text):
    """
    Cleans a single medical note text by removing boilerplate, placeholders, and other
    non-relevant content.
    """
    if not isinstance(text, str): # Handle potential non-string inputs
        return text
    else:
        # Remove boilerplate text like 'This is a sample medical note'
        text = re.sub(r'This is a sample medical note', '', text)
        # Remove placeholder text like 'N/A' or 'None'
        text = re.sub(r'N/A|None', '', text)
        # Remove punctuation and extra whitespace
        text = re.sub(r'[^\w\s]', '', text)
        text = re.sub(r'\s+', ' ', text)
        return text
```

```

    return ""
# Replace masked strings (like [**....**]) and placeholders (like __)
text = re.sub(r'\n', ' ', text)
text = re.sub(r'\[.*\.*?*\*\]', ' ', text)
text = re.sub(r'__+', ' ', text)

# Remove header sections by finding the start of the main content.
start_markers = [
    "History of Present Illness:",
    "Past Medical History:",
    "Social History:",
    "Family History:",
    "Physical Exam:",
    "Brief Hospital Course:",
    "Medications on Admission:",
    "Discharge Diagnosis:"
]

start_index = -1
for marker in start_markers:
    try:
        index = text.lower().index(marker.lower())
        if start_index == -1 or index < start_index:
            start_index = index
    except ValueError:
        continue

if start_index != -1:
    text = text[start_index:]

# Replace all excessive newlines, tabs, and multiple spaces with a single
text = text.replace('\t', ' ').replace('\r', ' ')
text = re.sub(r'\s+', ' ', text)

return text.strip()

def clean_discharge_notes(input_file: str, output_file: str):
    """
    Reads, cleans, and writes discharge summary notes in chunks.
    """
    print(f"Starting cleaning process for {input_file}...")

    # We will read and process the CSV in chunks to handle large files.
    # Adjust chunksize based on available RAM and file size
    chunk_size = 5000
    chunk_iter = pd.read_csv(input_file, chunksize=chunk_size, on_bad_lines=)

    first_chunk = True
    for chunk in chunk_iter:
        # Apply cleaning to the 'text' column
        chunk['text'] = chunk['text'].apply(clean_text)

        if first_chunk:
            chunk.to_csv(output_file, index=False, mode='w')
            first_chunk = False
        else:

```

```

        chunk.to_csv(output_file, index=False, mode='a', header=False)

    print(f"Processed a chunk of {chunk_size} rows and appended to {output_file}")

print(f"Cleaning complete. Cleaned data saved to {output_file}")

# Example usage (replace with your actual file paths)
# clean_discharge_notes('/content/sample_data/discharge.csv', '/content/proc...

```

Radiology

```
In [ ]: import re
import pandas as pd

def clean_text(text):
    """
    Cleans a single medical note text by removing boilerplate text, placeholders, and excessive whitespace.
    """
    if not isinstance(text, str): # Handle potential non-string inputs
        return ""
    # Replace masked strings (like [***...**]) and placeholders (like __)
    text = re.sub(r'\n', ' ', text)
    text = re.sub(r'\[*.*?*\*\]', ' ', text)
    text = re.sub(r'__+', ' ', text)

    # Remove header sections by finding the start of the main content.
    start_markers = [
        "EXAMINATION:",
        "INDICATION:",
        "TECHNIQUE:",
        "FINDINGS:",
        "IMPRESSION:"
    ]

    start_index = -1
    for marker in start_markers:
        try:
            index = text.lower().index(marker.lower())
            if start_index == -1 or index < start_index:
                start_index = index
        except ValueError:
            continue

    if start_index != -1:
        text = text[start_index:]

    # Replace all excessive newlines, tabs, and multiple spaces with a single space.
    text = text.replace('\t', ' ').replace('\r', ' ')
    text = re.sub(r'\s+', ' ', text)

    return text.strip()

def clean_radiology_notes(input_file: str, output_file: str):
    """
    Reads, cleans, and writes radiology notes in chunks.
    """

```

```

"""
print(f"Starting cleaning process for {input_file}...")

# We will read and process the CSV in chunks to handle large files.
# Adjust chunksize based on available RAM and file size
chunk_size = 5000
chunk_iter = pd.read_csv(input_file, chunksize=chunk_size, on_bad_lines=)

first_chunk = True
for chunk in chunk_iter:
    # Apply cleaning to the 'text' column
    chunk['text'] = chunk['text'].apply(clean_text)

    if first_chunk:
        chunk.to_csv(output_file, index=False, mode='w')
        first_chunk = False
    else:
        chunk.to_csv(output_file, index=False, mode='a', header=False)

    print(f"Processed a chunk of {chunk_size} rows and appended to {output_file}")

print(f"Cleaning complete. Cleaned data saved to {output_file}")

# Example usage (replace with your actual file paths)
# clean_radiology_notes('/content/sample_data/radiology.csv', '/content/proc

```

Combine Structured & Unstructured Data

```

In [ ]: import pandas as pd
import argparse # Import the argparse module
from pathlib import Path # Import Path

def merge_structured_unstructured_data(discharge_file: Path, radiology_file: Path):
    """
    merge_structured_unstructured_data function to read, merge, and write the data.
    Accepts file paths as arguments.
    """

    # Remove argparse logic as we are passing paths directly
    # parser = argparse.ArgumentParser(description='Join discharge, radiology, and features CSV files')
    # parser.add_argument('discharge_file', help='Path to the cleaned discharge CSV file')
    # parser.add_argument('radiology_file', help='Path to the cleaned radiology CSV file')
    # parser.add_argument('features_file', help='Path to the features CSV.')
    # parser.add_argument('output_file', help='Path to the output merged CSV file')
    # args = parser.parse_args()

    print("Reading CSV files...")
    try:
        discharge_df = pd.read_csv(discharge_file)
        radiology_df = pd.read_csv(radiology_file)
        features_df = pd.read_csv(features_file)
    except FileNotFoundError as e:

```

```

print(f"Error: {e.filename} not found.")
# Create empty dataframes if files are not found to prevent further
discharge_df = pd.DataFrame(columns=['subject_id', 'text'])
radiology_df = pd.DataFrame(columns=['subject_id', 'text'])
features_df = pd.DataFrame(columns=['subject_id'])
print("Using empty DataFrames due to missing input files.")

print("Aggregating notes by subject_id...")
# Aggregate notes by subject_id, joining all notes for a patient into a
# Ensure subject_id is treated as a common key
discharge_agg = discharge_df.groupby('subject_id')['text'].apply(' '.join)
discharge_agg.rename(columns={'text': 'discharge_text'}, inplace=True)

radiology_agg = radiology_df.groupby('subject_id')['text'].apply(' '.join)
radiology_agg.rename(columns={'text': 'radiology_text'}, inplace=True)

print("Merging dataframes on subject_id...")
# Merge the aggregated notes with the features dataframe.
# We use a left merge to keep all subjects from the features file.
merged_df = pd.merge(features_df, discharge_agg, on='subject_id', how='left')
final_df = pd.merge(merged_df, radiology_agg, on='subject_id', how='left')

# Fill NaN values in the new text columns with an empty string
final_df['discharge_text'] = final_df['discharge_text'].fillna('')
final_df['radiology_text'] = final_df['radiology_text'].fillna('')

print(f"Saving merged data to {output_file}...")
final_df.to_csv(output_file, index=False)
print(f"Merge complete. Final data saved to {output_file}")

# # Define file paths based on notebook structure
# # Assuming processed files are in the 'processed' directory
# processed_dir = Path("/content/processed") # Use the processed directory path
# discharge_input_path = processed_dir / "discharge_cleaned.csv"
# radiology_input_path = processed_dir / "radiology_cleaned.csv"
# # Note: The features file name might vary based on whether feature selection was used
# # Using features_full.csv as a default, but might need adjustment if selected features changed
# features_input_path = processed_dir / "features_full.csv" # Or "features_selection.csv"
# output_path = processed_dir / "merged_data.csv" # Define an output file name

# # Call merge_structured_unstructured_data with the defined paths
# merge_structured_unstructured_data(discharge_input_path, radiology_input_path)

```

Exploratory Data Analysis

In []:

```

import numpy as np
import pandas as pd

```

```
In [ ]: import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import ttest_ind

In [ ]: !ls

admissions.csv      icustays.csv    omr.csv      processed
d_icd_diagnoses.csv labevents.csv  patients.csv
```

```
In [ ]: from pathlib import Path
import pandas as pd

# Corrected file path
large_file_path = Path('/content/drive/MyDrive/MSML_602_FP/Dataset/workspace_data/processed/final_cleaned_data.csv')

sample_size = 50000

sampled_chunks = []

chunk_size = 50000

print(f"Reading and sampling from large file in chunks: {large_file_path}")
print(f"Target sample size: {sample_size}")

total_rows_read = 0
for chunk in pd.read_csv(large_file_path, chunksize=chunk_size, low_memory=False):
    total_rows_read += len(chunk)

    sample_fraction = min(1.0, (sample_size - sum(len(sc)) for sc in sampled_chunks))
    if sample_fraction <= 0:
        break

    sampled_chunk = chunk.sample(frac=sample_fraction, replace=False, random_state=42)
    sampled_chunks.append(sampled_chunk)
    print(f"Read {total_rows_read} rows, sampled {len(sampled_chunk)} from this chunk")

    if sum(len(sc) for sc in sampled_chunks) >= sample_size:
        break

project_df = pd.concat(sampled_chunks, ignore_index=True)

print(f"\nFinished sampling. Total rows in sampled DataFrame: {len(project_df)}
```

```
Reading and sampling from large file in chunks: /content/drive/MyDrive/MSML_602_FP/Dataset/workspace_data/processed/final_cleaned_data.csv
Target sample size: 50000
Read 50000 rows, sampled 50000 from this chunk. Total sampled so far: 50000
```

```
Finished sampling. Total rows in sampled DataFrame: 50000
```

```
In [ ]: features_selected_df = pd.read_csv('/content/drive/MyDrive/MSML_602_FP/Dataset/workspace_data/features_selected.csv')

In [ ]: eda_df = project_df.copy()
```

```
In [ ]: # 1. IMMEDIATE DATA QUALITY CHECKS
# Check text data completeness - CRITICAL for RAG
text_completeness = {
    'discharge_text': eda_df['discharge_text'].notna().sum(),
    'radiology_text': eda_df['radiology_text'].notna().sum(),
    'both_available': ((eda_df['discharge_text'].notna()) &
                       (eda_df['radiology_text'].notna())).sum()
}
print(f"\nText data availability: {text_completeness}")

# 2. Check text length distributions - short texts may be artifacts
eda_df['discharge_text_len'] = eda_df['discharge_text'].str.len()
eda_df['radiology_text_len'] = eda_df['radiology_text'].str.len()

print("\n")

print(eda_df[['discharge_text_len', 'radiology_text_len']].describe())

# Flag suspiciously short texts (< 100 chars)
short_discharge = (eda_df['discharge_text_len'] < 100).sum()
print(f"\nSuspiciously short discharge notes: {short_discharge}")

```

Text data availability: {'discharge_text': np.int64(38868), 'radiology_text': np.int64(41239), 'both_available': np.int64(37489)}

	discharge_text_len	radiology_text_len
count	3.886800e+04	41239.000000
mean	6.031893e+04	32774.408497
std	9.948833e+04	43446.177613
min	7.080000e+02	25.000000
25%	1.211600e+04	5793.000000
50%	2.811900e+04	16904.000000
75%	6.826100e+04	42155.000000
max	1.088152e+06	318705.000000

Suspiciously short discharge notes: 0

```
In [ ]: rows, columns = eda_df.shape
print(f"Number of rows: {rows}")
print(f"Number of columns: {columns}")

```

Number of rows: 50000
Number of columns: 29

```
In [ ]: eda_df.head()
```

```
Out[ ]:    subject_id    hadm_id    admittime    dischtime    admission_type    admit_provider_id    a
```

			2183-08-	2183-08-			
0	10636107	20812092	19 16:01:00	22 18:45:00	OBSERVATION ADMIT	P48CS5	E
1	10182665	22538295	2126-05- 29 11:30:00	2126-06- 01 13:17:00	SURGICAL SAME DAY ADMISSION	P8323I	
2	10003019	20962108	2176-01- 06 15:52:00	2176-01- 14 18:09:00	EW EMER.	P84UKK	E
3	10244511	26255794	2191-11- 29 03:09:00	2191-12- 02 15:50:00	EW EMER.	P42H7G	
4	10741731	21866879	2140-01- 26 09:15:00	2140-01- 26 15:38:00	SURGICAL SAME DAY ADMISSION	P60IOB	

5 rows × 29 columns

```
In [ ]: print("\n Column Types")
eda_df.dtypes
```

Column Types

Out[]:

	0
subject_id	int64
hadm_id	int64
admittime	object
dischtime	object
admission_type	object
admit_provider_id	object
admission_location	object
discharge_location	object
insurance	object
language	object
marital_status	object
race	object
edregtime	object
edouttime	object
hospital_expire_flag	int64
los_hours	float64
ed_wait_minutes	float64
ed_stay_minutes	float64
died_in_hospital	int64
gender	object
anchor_age	int64
anchor_year	int64
anchor_year_group	object
dod_available	int64
age_group	object
discharge_text	object
radiology_text	object
discharge_text_len	float64
radiology_text_len	float64

dtype: object

```
In [ ]: print(f"Unique Patients (subject_id): {eda_df['subject_id'].nunique():,}")
print(f"Unique Hospitalizations (hadm_id): {eda_df['hadm_id'].nunique():,}")
print(f"Unique admit_provider_id (admit_provider_id): {eda_df['admit_provider_id'].nunique():,}")
print(f"Total Rows: {len(eda_df):,}")
print(f"Average Rows per Patient: {len(eda_df) / eda_df['subject_id'].nunique():,.2f}")
```

```
Unique Patients (subject_id): 20,649
Unique Hospitalizations (hadm_id): 50,000
Unique admit_provider_id (admit_provider_id): 1,507
Total Rows: 50,000
Average Rows per Patient: 2.42
```

```
In [ ]: from pathlib import Path
import pandas as pd

processed_dir = Path("/content/processed")
merged_file_path = processed_dir / "merged_data.csv"

if merged_file_path.exists():
    print(f'{merged_file_path} ')
    try:

        df_check = pd.read_csv(merged_file_path, nrows=5, low_memory=False)
        print("\n ")
        display(df_check)
        print(f'\n {df_check.columns.tolist()}' )
    except Exception as e:
        print(f'\n error: {e}' )
else:
    print(f'{merged_file_path} ')
```

'/content/processed/merged_data.csv'

```
In [ ]: duplicates = eda_df.duplicated(subset=['subject_id', 'hadm_id']).sum()
print(f"Duplicate Hospital Stays: {duplicates}")
```

```
Duplicate Hospital Stays: 0
```

```
In [ ]: def categorize_columns(df): # Changed parameter name to df for clarity
    # Detect by name
    temporal_cols = [c for c in df.columns if any(x in c.lower() for x in ['date', 'time'])]
    id_cols = [c for c in df.columns if 'id' in c.lower()]
    text_cols = [c for c in df.columns if 'text' in c.lower()]

    # Detect by dtype
    object_cols = df.select_dtypes(include=['object']).columns.tolist()
    numeric_cols = df.select_dtypes(include=[np.number]).columns.tolist()

    # Detect categorical (object columns not text/id/time)
    categorical_cols = [
        c for c in object_cols
        if c not in text_cols + id_cols + temporal_cols
    ]

    # Detect numerical (numeric columns not id or binary)
    # Use the passed DataFrame 'df' instead of 'project_df'
    binary_cols = [
```

```

        c for c in numeric_cols
        if df[c].dropna().nunique() == 2
    ]

    numerical_cols = [
        c for c in numeric_cols
        if c not in id_cols + binary_cols
    ]

    return {
        'temporal_cols': temporal_cols,
        'id_cols': id_cols,
        'text_cols': text_cols,
        'categorical_cols': categorical_cols,
        'numerical_cols': numerical_cols,
        'binary_cols': binary_cols
    }

# Example usage:
col_categories = categorize_columns(eda_df) # Pass eda_df to the function

# Pretty print the results
for k, v in col_categories.items():
    print(f"{k}: {v}")

temporal_cols: ['admittime', 'dischtime', 'edregtime', 'edouttime']
id_cols: ['subject_id', 'hadm_id', 'admit_provider_id']
text_cols: ['discharge_text', 'radiology_text', 'discharge_text_len', 'radiology_text_len']
categorical_cols: ['admission_type', 'admission_location', 'discharge_location', 'insurance', 'language', 'marital_status', 'race', 'gender', 'anchor_year_group', 'age_group']
numerical_cols: ['los_hours', 'ed_wait_minutes', 'ed_stay_minutes', 'anchor_age', 'anchor_year', 'discharge_text_len', 'radiology_text_len']
binary_cols: ['hospital_expire_flag', 'died_in_hospital', 'dod_available']

```

In []: `eda_df.describe()`

	subject_id	hadm_id	hospital_expire_flag	los_hours	ed_wait_min
count	5.000000e+04	5.000000e+04	50000.000000	49984.000000	34779.000000
mean	1.047456e+07	2.498851e+07	0.022040	114.684835	341.348000
std	2.704533e+05	2.884753e+06	0.146815	185.462732	259.667000
min	1.000003e+07	2.000002e+07	0.000000	0.033333	-1411.000000
25%	1.024563e+07	2.250185e+07	0.000000	26.716667	184.000000
50%	1.048028e+07	2.499292e+07	0.000000	67.258333	282.000000
75%	1.070485e+07	2.748968e+07	0.000000	134.404167	420.000000
max	1.094027e+07	2.999972e+07	1.000000	12373.500000	4218.000000

In []: `eda_df.describe().T`

Out[]:

		count	mean	std	min	2
	subject_id	500000.0	1.047456e+07	2.704533e+05	1.000003e+07	1.024563e
	hadm_id	500000.0	2.498851e+07	2.884753e+06	2.000002e+07	2.250185e
	hospital_expire_flag	500000.0	2.204000e-02	1.468151e-01	0.000000e+00	0.000000e
	los_hours	49984.0	1.146848e+02	1.854627e+02	3.333333e-02	2.671667e
	ed_wait_minutes	34779.0	3.413489e+02	2.596678e+02	-1.411000e+03	1.840000e
	ed_stay_minutes	34779.0	6.565459e+02	6.188900e+02	0.000000e+00	3.130000e
	died_in_hospital	500000.0	2.204000e-02	1.468151e-01	0.000000e+00	0.000000e
	anchor_age	500000.0	5.697286e+01	1.909222e+01	1.800000e+01	4.300000e
	anchor_year	500000.0	2.152244e+03	2.378516e+01	2.110000e+03	2.132000e
	dod_available	500000.0	2.667800e-01	4.422808e-01	0.000000e+00	0.000000e
	discharge_text_len	38868.0	6.031893e+04	9.948833e+04	7.080000e+02	1.211600e
	radiology_text_len	41239.0	3.277441e+04	4.344618e+04	2.500000e+01	5.793000e

```
In [ ]: missing_stats = pd.DataFrame({
    'column': eda_df.columns,
    'missing_count': eda_df.isnull().sum().values,
    'missing_pct': (eda_df.isnull().sum() / len(eda_df)) * 100).values,
    'dtype': eda_df.dtypes.values
}).sort_values('missing_pct', ascending=False)
```

```
In [ ]: print(missing_stats[missing_stats['missing_pct'] > 0])
```

	column	missing_count	missing_pct	dtype
12	edregtime	15221	30.442	object
16	ed_wait_minutes	15221	30.442	float64
17	ed_stay_minutes	15221	30.442	float64
13	edouttime	15221	30.442	object
7	discharge_location	13811	27.622	object
27	discharge_text_len	11132	22.264	float64
25	discharge_text	11132	22.264	object
28	radiology_text_len	8761	17.522	float64
26	radiology_text	8761	17.522	object
8	insurance	839	1.678	object
15	los_hours	16	0.032	float64
5	admit_provider_id	1	0.002	object

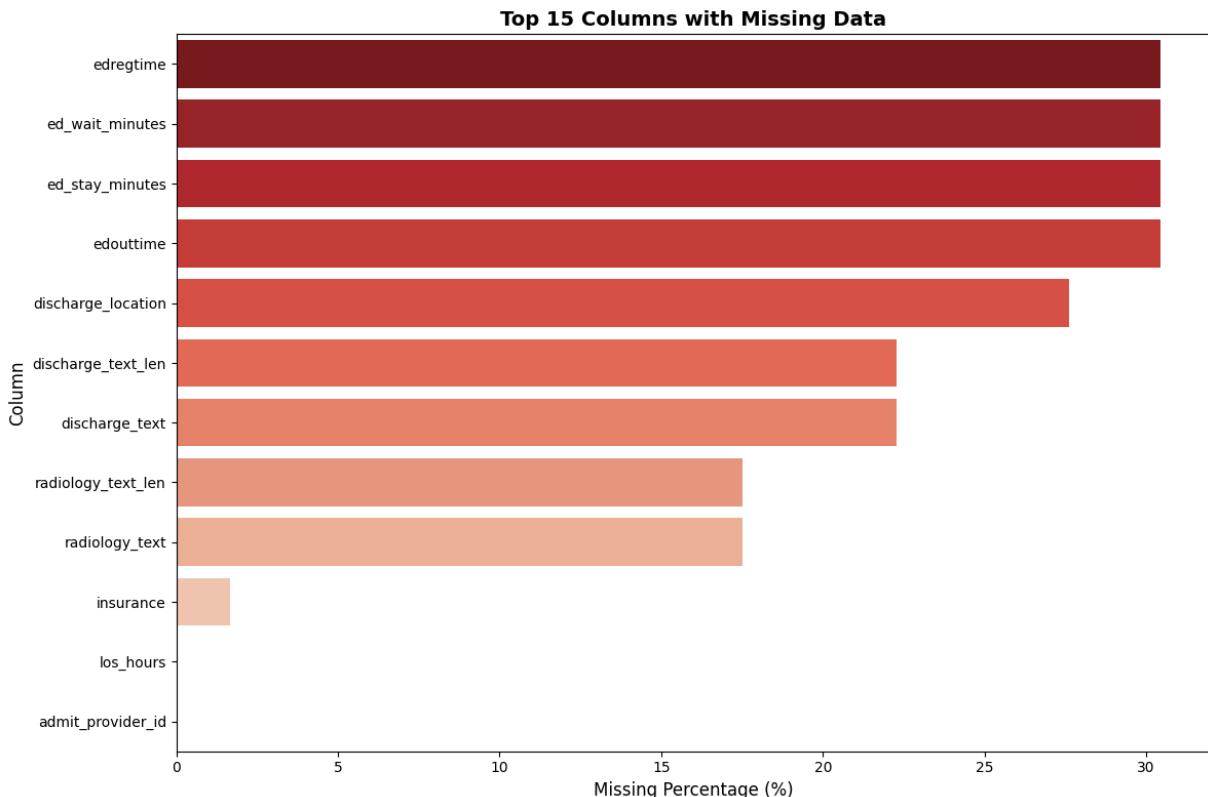
```
In [ ]: fig, ax = plt.subplots(figsize=(12, 8))
missing_plot_data = missing_stats[missing_stats['missing_pct'] > 0].head(15)
sns.barplot(data=missing_plot_data, x='missing_pct', y='column', palette='RdYlBu')
ax.set_xlabel('Missing Percentage (%)', fontsize=12)
ax.set_ylabel('Column', fontsize=12)
ax.set_title('Top 15 Columns with Missing Data', fontsize=14, fontweight='bold')
plt.tight_layout()
```

```
plt.savefig('missing_data_analysis.png', dpi=300, bbox_inches='tight')
plt.show()
```

```
/tmp/ipython-input-1429206783.py:3: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(data=missing_plot_data, x='missing_pct', y='column', palette
='Reds_r', ax=ax)
```



```
In [ ]: print("\nMissing Patterns by Context")

# ED columns missing pattern – expected for non-ED admissions
if 'ed_stay_minutes' in eda_df.columns:
    print("\nED Stay Minutes Missing by Admission Type:")
    ed_missing = eda_df.groupby('admission_type').agg({
        'ed_stay_minutes': lambda x: f"{x.isnull().sum()}/{len(x)} ({x.isnull().mean() * 100:.2f}%)"
    })
    print(ed_missing)

if 'ed_wait_minutes' in eda_df.columns:
    print("\nED Wait Minutes Missing by Admission Type:")
    ed_wait_missing = eda_df.groupby('admission_type').agg({
        'ed_wait_minutes': lambda x: f"{x.isnull().sum()}/{len(x)} ({x.isnull().mean() * 100:.2f}%)"
    })
    print(ed_wait_missing)

# Text data missingness – CRITICAL FOR RAG
print("\nText Data Completeness (CRITICAL FOR RAG)")
print(f"Discharge Text Available: {eda_df['discharge_text'].notna().sum():,}")
```

```
print(f"Radiology Text Available: {eda_df['radiology_text'].notna().sum():,}")
print(f"Both Texts Available: {((eda_df['discharge_text'].notna()) & (eda_df
```

Missing Patterns by Context

ED Stay Minutes Missing by Admission Type:
ed_stay_minutes

admission_type	
AMBULATORY OBSERVATION	600/613 (97.9%)
DIRECT EMER.	1759/2005 (87.7%)
DIRECT OBSERVATION	1657/2234 (74.2%)
ELECTIVE	1136/1147 (99.0%)
EU OBSERVATION	47/11103 (0.4%)
EW EMER.	689/16124 (4.3%)
OBSERVATION ADMIT	1166/7704 (15.1%)
SURGICAL SAME DAY ADMISSION	3966/3968 (99.9%)
URGENT	4201/5102 (82.3%)

ED Wait Minutes Missing by Admission Type:
ed_wait_minutes

admission_type	
AMBULATORY OBSERVATION	600/613 (97.9%)
DIRECT EMER.	1759/2005 (87.7%)
DIRECT OBSERVATION	1657/2234 (74.2%)
ELECTIVE	1136/1147 (99.0%)
EU OBSERVATION	47/11103 (0.4%)
EW EMER.	689/16124 (4.3%)
OBSERVATION ADMIT	1166/7704 (15.1%)
SURGICAL SAME DAY ADMISSION	3966/3968 (99.9%)
URGENT	4201/5102 (82.3%)

Text Data Completeness (CRITICAL FOR RAG)

Discharge Text Available: 38,868 (77.7%)

Radiology Text Available: 41,239 (82.5%)

Both Texts Available: 37,489 (75.0%)

```
In [ ]: eda_df['admittime'] = pd.to_datetime(eda_df['admittime'])
eda_df['dischtime'] = pd.to_datetime(eda_df['dischtime'])
if 'edregtime' in eda_df.columns:
    eda_df['edregtime'] = pd.to_datetime(eda_df['edregtime'])
if 'edouttime' in eda_df.columns:
    eda_df['edouttime'] = pd.to_datetime(eda_df['edouttime'])
```

```
In [ ]: if 'los_hours' not in eda_df.columns:
    eda_df['los_hours'] = (eda_df['dischtime'] - eda_df['admittime']).dt.total_seconds() / 3600
eda_df['los_days'] = eda_df['los_hours'] / 24

print("\nLength of Stay Statistics")
print("Hours:")
print(eda_df['los_hours'].describe())
print("\nDays:")
print(eda_df['los_days'].describe())
```

```
Length of Stay Statistics
Hours:
count    49984.000000
mean      114.684835
std       185.462732
min       0.033333
25%      26.716667
50%      67.258333
75%      134.404167
max      12373.500000
Name: los_hours, dtype: float64
```

```
Days:
count    49984.000000
mean      4.778535
std       7.727614
min       0.001389
25%      1.113194
50%      2.802431
75%      5.600174
max      515.562500
Name: los_days, dtype: float64
```

```
In [ ]: print("\nTemporal Data Quality Checks")
print(f"Negative LOS: {(eda_df['los_days'] < 0).sum()}")
print(f"Zero LOS: {(eda_df['los_days'] == 0).sum()}")
print(f"Same-day discharge (< 1 day): {(eda_df['los_days'] < 1).sum()}")
print(f"Extended stay (> 30 days): {(eda_df['los_days'] > 30).sum()}")
print(f"Very long stay (> 365 days): {(eda_df['los_days'] > 365).sum()}"
```

```
Temporal Data Quality Checks
Negative LOS: 0
Zero LOS: 0
Same-day discharge (< 1 day): 11242
Extended stay (> 30 days): 689
Very long stay (> 365 days): 1
```

```
In [ ]: if 'ed_wait_minutes' in eda_df.columns and eda_df['ed_wait_minutes'].notna():
    print("\nEmergency Department Timing")
    print("ED Wait Time (minutes):")
    print(eda_df['ed_wait_minutes'].describe())
    print("\nED Total Stay Time (minutes):")
    print(eda_df['ed_stay_minutes'].describe())
```

```
Emergency Department Timing
ED Wait Time (minutes):
count    34779.00000
mean     341.348946
std      259.667793
min     -1411.00000
25%      184.00000
50%      282.00000
75%      420.00000
max     4218.00000
Name: ed_wait_minutes, dtype: float64
```

```
ED Total Stay Time (minutes):
count    34779.00000
mean     656.545875
std      618.890005
min      0.00000
25%     313.00000
50%     466.00000
75%     776.00000
max    11683.00000
Name: ed_stay_minutes, dtype: float64
```

```
In [ ]: eda_df['admit_year'] = eda_df['admittime'].dt.year
eda_df['admit_month'] = eda_df['admittime'].dt.month
eda_df['admit_hour'] = eda_df['admittime'].dt.hour
eda_df['admit_dow'] = eda_df['admittime'].dt.dayofweek # 0=Monday
eda_df['admit_day_name'] = eda_df['admittime'].dt.day_name()
```

```
In [ ]: # Temporal Features
eda_df['admit_year'] = eda_df['admittime'].dt.year
eda_df['admit_month'] = eda_df['admittime'].dt.month
eda_df['admit_hour'] = eda_df['admittime'].dt.hour
eda_df['admit_dow'] = eda_df['admittime'].dt.dayofweek # 0=Monday
eda_df['admit_day_name'] = eda_df['admittime'].dt.day_name()

# Visualize Temporal Patterns
fig, axes = plt.subplots(2, 3, figsize=(18, 10))

# Admissions by year
eda_df['admit_year'].value_counts().sort_index().plot(kind='bar', ax=axes[0, 0])
axes[0, 0].set_title('Admissions by Year (Deidentified)', fontweight='bold')
axes[0, 0].set_xlabel('Year')
axes[0, 0].set_ylabel('Count')

# Admissions by month
month_counts = eda_df['admit_month'].value_counts().sort_index()
axes[0, 1].bar(range(1, 13), month_counts.values, color='coral')
axes[0, 1].set_xticks(range(1, 13))
axes[0, 1].set_xticklabels(['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun',
                           'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec'])
axes[0, 1].set_title('Admissions by Month', fontweight='bold')
axes[0, 1].set_xlabel('Month')
axes[0, 1].set_ylabel('Count')

# Admissions by hour
```

```

eda_df['admit_hour'].value_counts().sort_index().plot(kind='bar', ax=axes[0],
axes[0,2].set_title('Admissions by Hour of Day', fontweight='bold')
axes[0,2].set_xlabel('Hour')
axes[0,2].set_ylabel('Count')

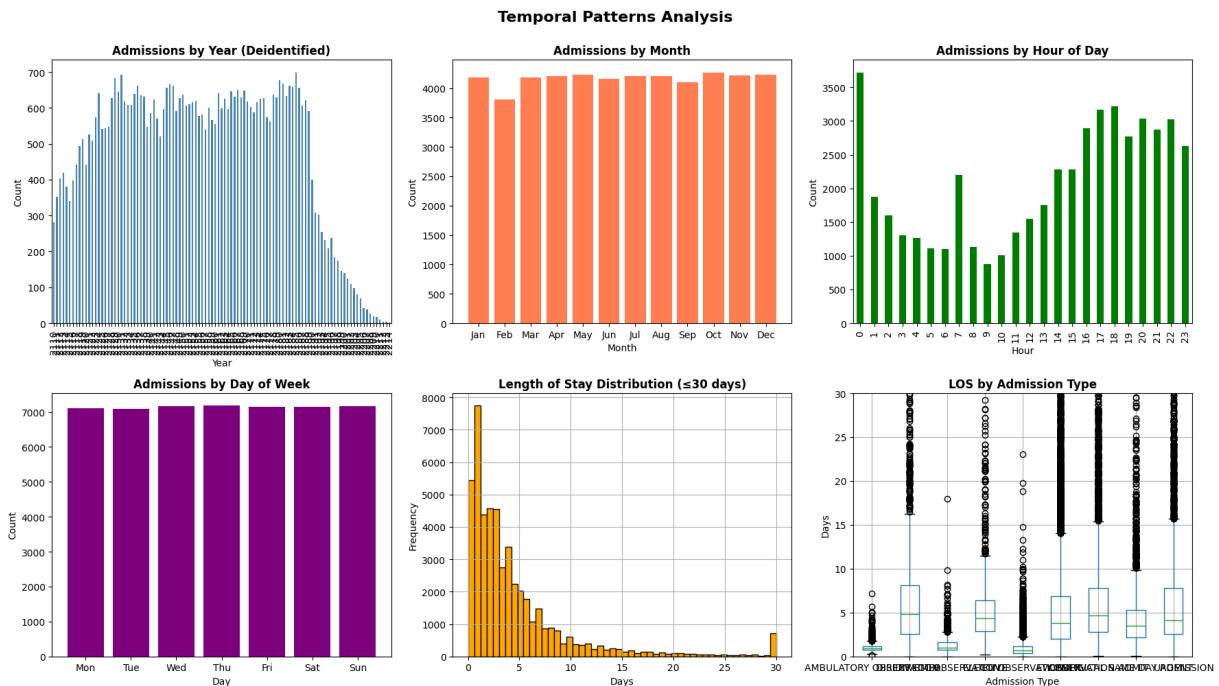
# Admissions by day of week
day_order = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']
dow_data = eda_df['admit_day_name'].value_counts().reindex(day_order)
axes[1,0].bar(range(7), dow_data.values, color='purple')
axes[1,0].set_xticks(range(7))
axes[1,0].set_xticklabels(['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun'])
axes[1,0].set_title('Admissions by Day of Week', fontweight='bold')
axes[1,0].set_xlabel('Day')
axes[1,0].set_ylabel('Count')

# LOS distribution (capped)
eda_df['los_days'].clip(upper=30).hist(bins=50, ax=axes[1,1], color='orange')
axes[1,1].set_title('Length of Stay Distribution ( $\leq 30$  days)', fontweight='bold')
axes[1,1].set_xlabel('Days')
axes[1,1].set_ylabel('Frequency')

# LOS by admission type
eda_df.boxplot(column='los_days', by='admission_type', ax=axes[1,2])
axes[1,2].set_title('LOS by Admission Type', fontweight='bold')
axes[1,2].set_xlabel('Admission Type')
axes[1,2].set_ylabel('Days')
axes[1,2].set_ylim(0, 30)
plt.suptitle('') # Remove default title

plt.suptitle('Temporal Patterns Analysis', fontsize=16, fontweight='bold', y=0.95)
plt.tight_layout()
plt.savefig('temporal_analysis.png', dpi=300, bbox_inches='tight')
plt.show()

```



```
In [ ]: # Anchor Year Analysis
print("\nAnchor Year Groups")
print(eda_df['anchor_year_group'].value_counts().sort_index())
```

```
Anchor Year Groups
anchor_year_group
2008 - 2010    20956
2011 - 2013    10144
2014 - 2016     8484
2017 - 2019     6537
2020 - 2022     3879
Name: count, dtype: int64
```

```
In [ ]: print("\nAge Distribution")
eda_df['anchor_age'].describe()
```

```
Age Distribution
```

```
Out[ ]:      anchor_age
```

	anchor_age
count	50000.000000
mean	56.972860
std	19.092218
min	18.000000
25%	43.000000
50%	58.000000
75%	72.000000
max	91.000000

```
dtype: float64
```

```
In [ ]: print(f"\nAge-censored patients (age=91): {(eda_df['anchor_age'] == 91).sum()}")
Age-censored patients (age=91): 1,428 (2.9%)
```

```
In [ ]: print("\nAge Groups")
print(eda_df['age_group'].value_counts().sort_index())
```

```
Age Groups
age_group
40-59      15450
60-74      13765
75+        10205
<40        10580
Name: count, dtype: int64
```

```
In [ ]: print("\nGender Distribution")
gender_counts = eda_df['gender'].value_counts()
print(gender_counts)
```

```
print("\nPercentages:")
print(eda_df['gender'].value_counts(normalize=True) * 100)
```

Gender Distribution
gender
F 25596
M 24404
Name: count, dtype: int64

Percentages:
gender
F 51.192
M 48.808
Name: proportion, dtype: float64

```
In [ ]: print("\nRace/Ethnicity Distribution")
print(eda_df['race'].value_counts().head(10))
```

Race/Ethnicity Distribution
race
WHITE 30889
BLACK/AFRICAN AMERICAN 6948
OTHER 1864
UNKNOWN 1362
WHITE - OTHER EUROPEAN 1234
HISPANIC/LATINO - PUERTO RICAN 945
ASIAN 737
HISPANIC OR LATINO 684
ASIAN - CHINESE 679
BLACK/CAPE VERDEAN 592
Name: count, dtype: int64

```
In [ ]: print("\nMarital Status Distribution")
print(eda_df['marital_status'].value_counts())
```

Marital Status Distribution
marital_status
MARRIED 20909
SINGLE 18648
WIDOWED 5278
DIVORCED 3915
UNKNOWN 1250
Name: count, dtype: int64

```
In [ ]: fig, axes = plt.subplots(2, 3, figsize=(18, 10))

# Age histogram (excluding censored)
eda_df[eda_df['anchor_age'] < 91]['anchor_age'].hist(bins=40, ax=axes[0,0],
                                                       color='skyblue', edgecolor='black')
axes[0,0].set_title('Age Distribution (excluding censored)', fontweight='bold')
axes[0,0].set_xlabel('Age')
axes[0,0].set_ylabel('Frequency')

# Gender distribution
gender_counts.plot(kind='pie', ax=axes[0,1], autopct='%1.1f%%', colors=['lightblue', 'pink'])
axes[0,1].set_title('Gender Distribution', fontweight='bold')
axes[0,1].set_ylabel('')
```

```

# Age groups
eda_df['age_group'].value_counts().sort_index().plot(kind='bar', ax=axes[0,2])
axes[0,2].set_title('Age Groups', fontweight='bold')
axes[0,2].set_xlabel('Age Group')
axes[0,2].set_ylabel('Count')
axes[0,2].tick_params(axis='x', rotation=45)

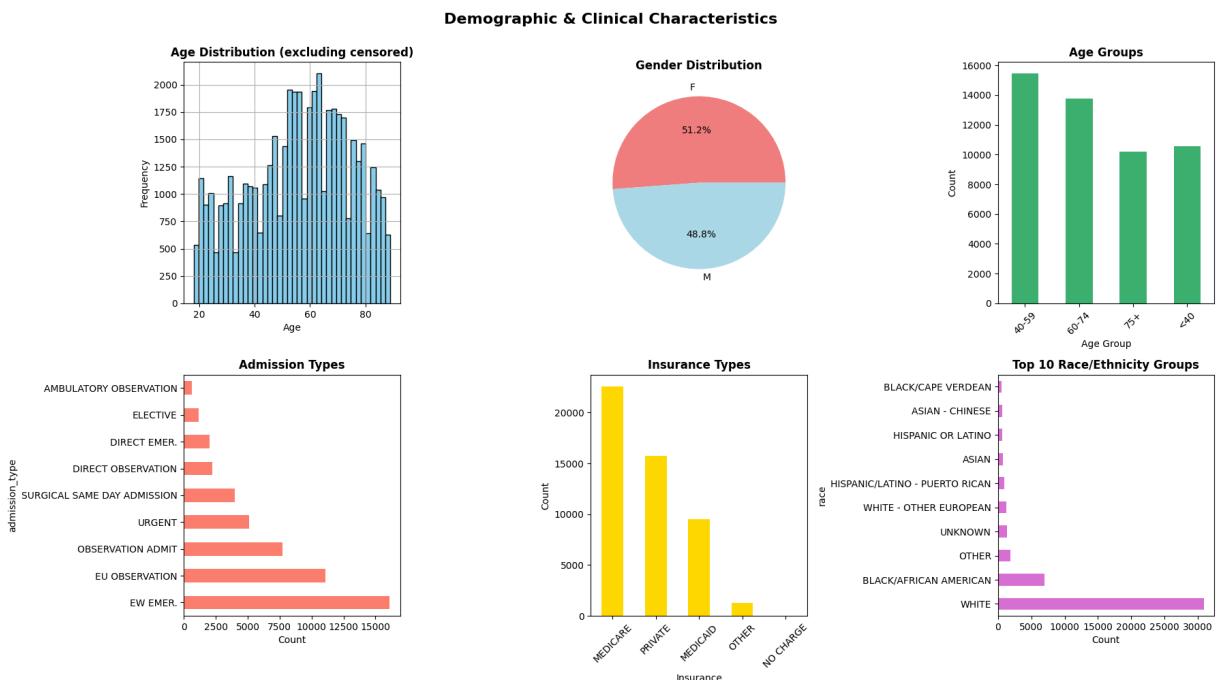
# Admission type
eda_df['admission_type'].value_counts().plot(kind='barh', ax=axes[1,0], color='red')
axes[1,0].set_title('Admission Types', fontweight='bold')
axes[1,0].set_xlabel('Count')

# Insurance type
eda_df['insurance'].value_counts().plot(kind='bar', ax=axes[1,1], color='gold')
axes[1,1].set_title('Insurance Types', fontweight='bold')
axes[1,1].set_xlabel('Insurance')
axes[1,1].set_ylabel('Count')
axes[1,1].tick_params(axis='x', rotation=45)

# Top race/ethnicity
eda_df['race'].value_counts().head(10).plot(kind='barh', ax=axes[1,2], color='teal')
axes[1,2].set_title('Top 10 Race/Ethnicity Groups', fontweight='bold')
axes[1,2].set_xlabel('Count')

plt.suptitle('Demographic & Clinical Characteristics', fontsize=16, fontweight='bold')
plt.tight_layout()
plt.savefig('demographics_analysis.png', dpi=300, bbox_inches='tight')
plt.show()

```



```
In [ ]: print("\nAdmission Characteristics")
print("\nAdmission Types:")
print(eda_df['admission_type'].value_counts())
```

Admission Characteristics

Admission Types:

admission_type	
EW EMER.	16124
EU OBSERVATION	11103
OBSERVATION ADMIT	7704
URGENT	5102
SURGICAL SAME DAY ADMISSION	3968
DIRECT OBSERVATION	2234
DIRECT EMER.	2005
ELECTIVE	1147
AMBULATORY OBSERVATION	613
Name: count, dtype: int64	

```
In [ ]: print("\nTop 10 Admission Locations:")
print(eda_df['admission_location'].value_counts().head(10))
```

Top 10 Admission Locations:

admission_location	
EMERGENCY ROOM	22509
PHYSICIAN REFERRAL	14824
TRANSFER FROM HOSPITAL	5247
WALK-IN/SELF REFERRAL	3813
CLINIC REFERRAL	1160
PROCEDURE SITE	758
INTERNAL TRANSFER TO OR FROM PSYCH	576
TRANSFER FROM SKILLED NURSING FACILITY	553
PACU	497
INFORMATION NOT AVAILABLE	40
Name: count, dtype: int64	

```
In [ ]: print("\nTop 10 Discharge Locations:")
print(eda_df['discharge_location'].value_counts().head(10))
```

Top 10 Discharge Locations:

discharge_location	
HOME	17830
HOME HEALTH CARE	8926
SKILLED NURSING FACILITY	4769
REHAB	1326
DIED	1099
CHRONIC/LONG TERM ACUTE CARE	730
HOSPICE	481
AGAINST ADVICE	314
PSYCH FACILITY	274
ACUTE HOSPITAL	223
Name: count, dtype: int64	

```
In [ ]: print("\nLanguage Distribution:")
print(eda_df['language'].value_counts().head(10))
```

```
Language Distribution:  
language  
ENGLISH           44767  
SPANISH            1833  
CHINESE             755  
RUSSIAN             657  
KABUVERDIANU        480  
PORTUGUESE          306  
HAITIAN              205  
OTHER                 162  
VIETNAMESE           133  
MODERN GREEK (1453-)    116  
Name: count, dtype: int64
```

```
In [ ]: if 'hospital_expire_flag' in eda_df.columns:  
    mortality_col = 'hospital_expire_flag'  
else:  
    mortality_col = 'died_in_hospital'  
  
mortality_rate = eda_df[mortality_col].mean() * 100  
n_deaths = eda_df[mortality_col].sum()  
  
print(f"\nOverall Mortality")  
print(f"In-hospital Deaths: {n_deaths:,}")  
print(f"In-hospital Mortality Rate: {mortality_rate:.2f}%")
```

```
Overall Mortality  
In-hospital Deaths: 1,102  
In-hospital Mortality Rate: 2.20%
```

```
In [ ]: if 'hospital_expire_flag' in eda_df.columns and 'died_in_hospital' in eda_df:  
    print("\nMortality Flag Comparison")  
    print(pd.crosstab(eda_df['hospital_expire_flag'], eda_df['died_in_hospital'],  
                      rownames=['hospital_expire_flag'], colnames=['died_in_hospital']))  
  
# DOD availability  
if 'dod_available' in eda_df.columns:  
    dod_available = eda_df['dod_available'].sum()  
    print(f"\nDate of Death Available: {dod_available:,} ({dod_available/len(eda_df)*100:.2f}%")
```

```
Mortality Flag Comparison  
died_in_hospital      0      1  
hospital_expire_flag  
0                  48898      0  
1                  0     1102
```

```
Date of Death Available: 13,339 (26.68%)
```

```
In [ ]: print("\nMortality by Age Group")  
mortality_by_age = eda_df.groupby('age_group')[mortality_col].agg(['sum', 'count'])  
mortality_by_age['rate_%'] = mortality_by_age['mean'] * 100  
print(mortality_by_age)
```

```
Mortality by Age Group
      sum   count      mean    rate_%
age_group
40-59       200  15450  0.012945  1.294498
60-74       376  13765  0.027316  2.731566
75+        487  10205  0.047722  4.772171
<40        39   10580  0.003686  0.368620
```

```
In [ ]: print("\nMortality by Gender")
mortality_by_gender = eda_df.groupby('gender')[mortality_col].agg(['sum', 'c
mortality_by_gender['rate_%'] = mortality_by_gender['mean'] * 100
print(mortality_by_gender)
```

```
Mortality by Gender
      sum   count      mean    rate_%
gender
F        487  25596  0.019026  1.902641
M        615  24404  0.025201  2.520079
```

```
In [ ]: print("\nMortality by Race/Ethnicity (Top 10)")
mortality_by_race = eda_df.groupby('race')[mortality_col].agg(['sum', 'count
mortality_by_race['rate_%'] = mortality_by_race['mean'] * 100
print(mortality_by_race.sort_values('count', ascending=False).head(10))
```

```
Mortality by Race/Ethnicity (Top 10)
      sum   count      mean    rate_%
race
WHITE                  613  30889  0.019845  1.984525
BLACK/AFRICAN AMERICAN  96   6948   0.013817  1.381693
OTHER                  34   1864   0.018240  1.824034
UNKNOWN                 162  1362   0.118943  11.894273
WHITE - OTHER EUROPEAN   27   1234   0.021880  2.188006
HISpanic/LATINO - PUERTO RICAN  7   945    0.007407  0.740741
ASIAN                   19   737    0.025780  2.578019
HISpanic OR LATINO      6   684    0.008772  0.877193
ASIAN - CHINESE          16   679    0.023564  2.356406
BLACK/CAPE VERDEAN       7   592    0.011824  1.182432
```

```
In [ ]: print("\nMortality by Admission Type")
mortality_by_admission = eda_df.groupby('admission_type')[mortality_col].agg(
mortality_by_admission['rate_%'] = mortality_by_admission['mean'] * 100
print(mortality_by_admission.sort_values('rate_%', ascending=False))
```

```
Mortality by Admission Type
      sum   count      mean    rate_%
admission_type
URGENT                  236  5102   0.046256  4.625637
EW EMER.                 627  16124   0.038886  3.888613
DIRECT EMER.              55   2005   0.027431  2.743142
OBSERVATION ADMIT         159  7704   0.020639  2.063863
ELECTIVE                  6   1147   0.005231  0.523104
SURGICAL SAME DAY ADMISSION  8   3968   0.002016  0.201613
DIRECT OBSERVATION         2   2234   0.000895  0.089526
EU OBSERVATION              9  11103   0.000811  0.081059
AMBULATORY OBSERVATION        0   613    0.000000  0.000000
```

```
In [ ]: print("\nMortality by Insurance Type")
mortality_by_insurance = eda_df.groupby('insurance')[mortality_col].agg(['sum', 'count', 'mean'])
mortality_by_insurance['rate_%'] = mortality_by_insurance['mean'] * 100
print(mortality_by_insurance.sort_values('rate_%', ascending=False))
```

```
Mortality by Insurance Type
      sum   count      mean    rate_%
insurance
OTHER      52    1271  0.040913  4.091267
MEDICARE    703   22595  0.031113  3.111308
PRIVATE     204   15754  0.012949  1.294909
MEDICAID     99    9499  0.010422  1.042215
NO CHARGE     0     42  0.000000  0.000000
```

```
In [ ]: print("\nData Validation: Mortality vs Discharge Location")
mortality_discharge_check = pd.crosstab(
    eda_df[mortality_col],
    eda_df['discharge_location'],
    margins=True,
    margins_name='Total'
)
print(mortality_discharge_check)
```

```
Data Validation: Mortality vs Discharge Location
discharge_location    ACUTE HOSPITAL  AGAINST ADVICE  ASSISTED LIVING  \
hospital_expire_flag
0                      223           314            63
1                      0             0              0
Total                  223           314            63

discharge_location    CHRONIC/LONG TERM ACUTE CARE  DIED  HEALTHCARE FACILIT \
Y \
hospital_expire_flag
0                      730           24
7
1                      0           1075
0
Total                  730           1099
7

discharge_location    HOME  HOME HEALTH CARE  HOSPICE  OTHER FACILITY  \
hospital_expire_flag
0                     17819          8926        481         147
1                      11            0           0            0
Total                  17830          8926        481         147

discharge_location    PSYCH FACILITY  REHAB  SKILLED NURSING FACILITY  Total
hospital_expire_flag
0                      274          1325
1                      0            1
Total                  274          1326
4766  35099
3       1090
4769  36189
```

```
In [ ]: fig, axes = plt.subplots(2, 2, figsize=(15, 10))

# Mortality by age group
mortality_by_age['rate_%'].plot(kind='bar', ax=axes[0,0], color='darkred')
```

```

axes[0,0].set_title('Mortality Rate by Age Group', fontweight='bold')
axes[0,0].set_xlabel('Age Group')
axes[0,0].set_ylabel('Mortality Rate (%)')
axes[0,0].tick_params(axis='x', rotation=45)

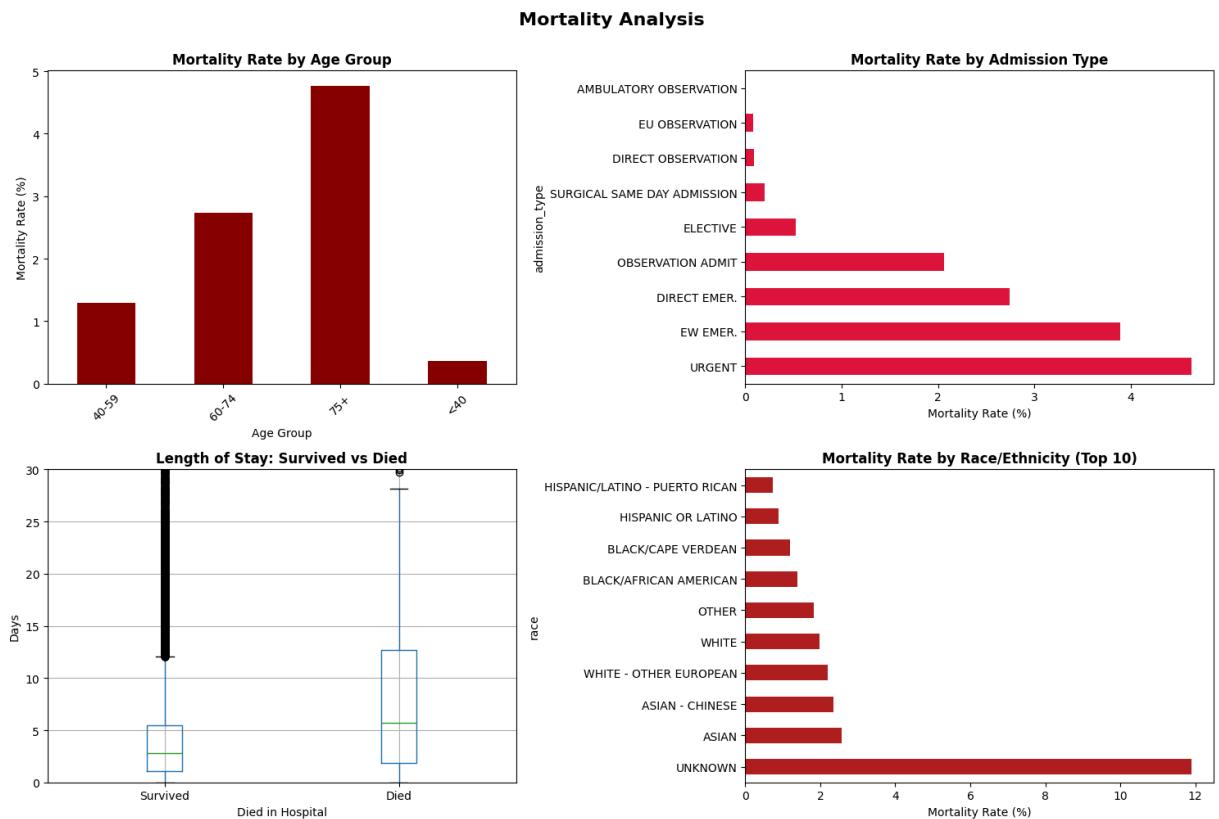
# Mortality by admission type
mortality_by_admission['rate_%'].sort_values(ascending=False).plot(kind='bar')
axes[0,1].set_title('Mortality Rate by Admission Type', fontweight='bold')
axes[0,1].set_xlabel('Mortality Rate (%)')

# LOS comparison: survived vs died
eda_df.boxplot(column='los_days', by=mortality_col, ax=axes[1,0])
axes[1,0].set_title('Length of Stay: Survived vs Died', fontweight='bold')
axes[1,0].set_xlabel('Died in Hospital')
axes[1,0].set_ylabel('Days')
axes[1,0].set_xticklabels(['Survived', 'Died'])
axes[1,0].set_ylim(0, 30)
plt.suptitle('') # Remove default title

# Mortality by race (top 10)
top_races = eda_df['race'].value_counts().head(10).index
mortality_by_race_top = eda_df[eda_df['race'].isin(top_races)].groupby('race')
mortality_by_race_top.sort_values(ascending=False).plot(kind='barh', ax=axes[1,1])
axes[1,1].set_title('Mortality Rate by Race/Ethnicity (Top 10)', fontweight='bold')
axes[1,1].set_xlabel('Mortality Rate (%)')

plt.suptitle('Mortality Analysis', fontsize=16, fontweight='bold', y=1.00)
plt.tight_layout()
plt.savefig('mortality_analysis.png', dpi=300, bbox_inches='tight')
plt.show()

```



```
In [ ]: def analyze_clinical_text(text_series, text_len_series, name):  
    """Comprehensive text analysis for RAG preparation"""  
  
    print(f"\n{'='*70}")  
    print(f"{name.upper()} ANALYSIS")  
    print(f"{'='*70}")  
  
    # Remove nulls  
    texts = text_series.dropna()  
    print(f"\nTotal Available: {len(texts)}, {(len(texts)/len(eda_df))*100:.2f}%")  
  
    if len(texts) == 0:  
        print("No text data available!")  
        return  
  
    # Character length analysis (use pre-calculated if available)  
    if text_len_series is not None:  
        char_lengths = text_len_series.dropna()  
    else:  
        char_lengths = texts.str.len()  
  
    print(f"\nCharacter Length Statistics")  
    print(char_lengths.describe())  
  
    # Word count analysis  
    word_counts = texts.str.split().str.len()  
    print(f"\nWord Count Statistics")  
    print(word_counts.describe())  
  
    # Identify suspiciously short texts  
    short_threshold = 100  
    very_short = (char_lengths < short_threshold).sum()  
    print(f"\nSuspiciously Short Texts (< {short_threshold} chars): {very_short}")  
  
    # Medical terminology presence  
    medical_terms = {  
        'diagnosis': r'\bdiagnos(is|es|ed|ing)\b',  
        'history': r'\bhistory\b',  
        'examination': r'\bexamin(ation|ed|ing)\b',  
        'treatment': r'\btreat(ment|ed|ing)\b',  
        'symptoms': r'\bsymptom(s)?\b',  
        'patient': r'\bpatient(s)?\b',  
        'admitted': r'\badmit(ted|ting)?\b',  
        'procedure': r'\bprocedure(s)?\b',  
        'medication': r'\bmedication(s)?\b',  
        'discharge': r'\bdischarge(d)?\b'  
    }  
  
    print(f"\nMedical Terminology Presence")  
    for term, pattern in medical_terms.items():  
        count = texts.str.contains(pattern, case=False, regex=True, na=False)  
        print(f"{term.capitalize():15}: {count:6}, {(count/len(texts))*100:.5f}%")  
  
    # Visualize distributions  
    fig, axes = plt.subplots(1, 2, figsize=(14, 5))
```

```

# Character length distribution
char_lengths.clip(upper=char_lengths.quantile(0.95)).hist(bins=50, ax=axes[0], color='teal')
axes[0].set_title(f'{name} - Character Length Distribution (95th percentile)')
axes[0].set_xlabel('Characters')
axes[0].set_ylabel('Frequency')
axes[0].axvline(char_lengths.median(), color='red', linestyle='--', label='Median')
axes[0].legend()

# Word count distribution
word_counts.clip(upper=word_counts.quantile(0.95)).hist(bins=50, ax=axes[1], color='darkorange')
axes[1].set_title(f'{name} - Word Count Distribution (95th percentile)')
axes[1].set_xlabel('Words')
axes[1].set_ylabel('Frequency')
axes[1].axvline(word_counts.median(), color='red', linestyle='--', label='Median')
axes[1].legend()

plt.tight_layout()
plt.savefig(f'{name.lower().replace(" ", "_")}_analysis.png', dpi=300, bbox_inches='tight')
plt.show()

# Sample text
print(f"\nSample Text (First 500 characters)")
sample_text = texts.iloc[0][:500]
print(sample_text)
print("...")

```

In []:

```

#####
analyze_clinical_text(eda_df['discharge_text'],
                      eda_df['discharge_text_len'] if 'discharge_text_len' in eda_df.columns
                      else None, "DISCHARGE NOTES")
analyze_clinical_text(eda_df['radiology_text'],
                      eda_df['radiology_text_len'] if 'radiology_text_len' in eda_df.columns
                      else None, "RADIOLOGY REPORTS")
#####

```

Out[]:

```

'\nanalyze_clinical_text(eda_df['discharge_text'],\neda_df['discharge_text_len'] if 'discharge_text_len' in eda_df.columns\nelse None, "DISCHARGE NOTES")\nanalyze_clinical_text(\neda_df['radiology_text'],\neda_df['radiology_text_len'] if 'radiology_text_len' in eda_df.columns\nelse None, "RADIOLOGY REPORTS")\n'

```

Computing and visualizing correlations among key numerical and binary variables to identify relationships with in-hospital mortality. This helps highlight which features are most strongly associated with patient outcomes.

In []:

```

temporal_cols = ['admittime', 'dischtime', 'edregtime', 'edouttime']
id_cols = ['subject_id', 'hadm_id', 'admit_provider_id']
text_cols = ['discharge_text', 'radiology_text', 'discharge_text_len', 'radiology_text_len']
categorical_cols = ['admission_type', 'admission_location', 'discharge_location',
                    'language', 'marital_status', 'race', 'gender', 'anchor_hospital']
numerical_cols = ['los_hours', 'ed_wait_minutes', 'ed_stay_minutes', 'anchor_hospital']

```

```
'discharge_text_len', 'radiology_text_len']
binary_cols = ['hospital_expire_flag', 'died_in_hospital', 'dod_available']

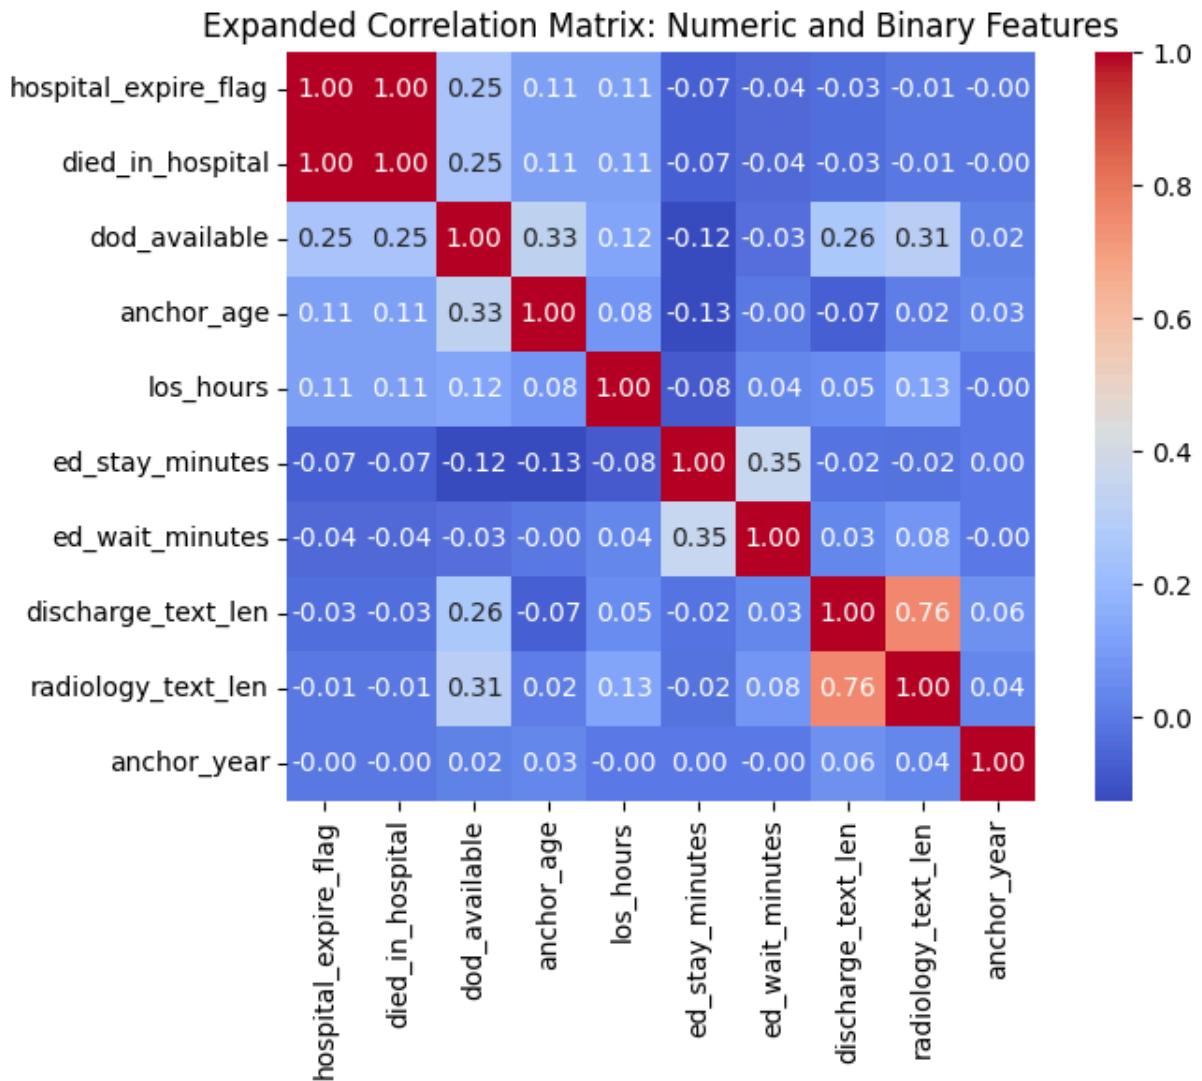
# Subset and prepare correlation-ready DataFrame
corr_cols = numerical_cols + binary_cols
corr_df = eda_df[corr_cols].copy()

# Ensure binary columns are numeric (just in case)
for c in binary_cols:
    if corr_df[c].dtype == 'object':
        corr_df[c] = corr_df[c].map({'Yes': 1, 'No': 0}).astype(float)

# Compute correlation matrix
corr_matrix = corr_df.corr(numeric_only=True)

# Sort by strength of correlation with mortality
if 'died_in_hospital' in corr_matrix.columns:
    sorted_cols = corr_matrix['died_in_hospital'].abs().sort_values(ascending=False)
    corr_matrix = corr_matrix.loc[sorted_cols, sorted_cols]

# Visualization
plt.figure(figsize=(8, 6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f", square=True)
plt.title("Expanded Correlation Matrix: Numeric and Binary Features")
plt.tight_layout()
plt.show()
```



Identifying key factors potentially linked to patient outcomes.

```
In [ ]: # Top correlated features with mortality
top_corr = corr_matrix['died_in_hospital'].drop('died_in_hospital').sort_values
print("\nTop correlations with mortality (died_in_hospital):")
print(top_corr)
```

Top correlations with mortality (died_in_hospital):

hospital_expire_flag	1.000000
dod_available	0.248878
anchor_age	0.106551
los_hours	0.106355
ed_stay_minutes	-0.068048
ed_wait_minutes	-0.044221
discharge_text_len	-0.031312
radiology_text_len	-0.008705
anchor_year	-0.003191

Name: died_in_hospital, dtype: float64

The strongest correlation with in-hospital mortality is the hospital_expire_flag as expected, this indicates perfect alignment between the two indicators of death. Other moderate correlations include date of death availability (dod_available) and length of

stay (los_hours), suggesting that longer hospital stays and older age may be associated with higher mortality.

Most other features show weak correlations, indicating limited linear relationships with mortality.

Statistical Analysis: Length of Stay and Mortality (Welch's t-test)

Performing a Welch's t-test to compare the average length of stay between patients who died in the hospital and those who survived.

This is to test whether the difference in means (averages) is statistically significant.

Hypotheses:

- **Null Hypothesis (H0):** The average length of hospital stay for patients who died is equal to the average length of hospital stay for patients who survived.
- **Alternative Hypothesis (H1):** The average length of hospital stay for patients who died is not equal to the average length of hospital stay for patients who survived.

```
In [ ]: # Split groups
died = eda_df.loc[eda_df['died_in_hospital'] == 1, 'los_hours'].dropna()
survived = eda_df.loc[eda_df['died_in_hospital'] == 0, 'los_hours'].dropna()

# Performing Welch's t-test
t_stat, p_val = ttest_ind(died, survived, equal_var=False)

print(f"T-statistic: {t_stat:.3f}")
print(f"P-value: {p_val:.4f}")

if p_val < 0.05:
    print("Statistically significant difference: hospital stay is longer for")
else:
    print("No statistically significant difference detected.")

T-statistic: 11.435
P-value: 0.0000
Statistically significant difference: hospital stay is longer for patients w
ho died.
```

The results show a highly significant difference in hospital stay length between patients who died and those who survived. Patients who died tended to have **substantially longer hospital stays**, as indicated by the **large t-statistic and zero p-value**.

This visualization compares the distribution of hospital stay lengths between patients who survived and those who died.

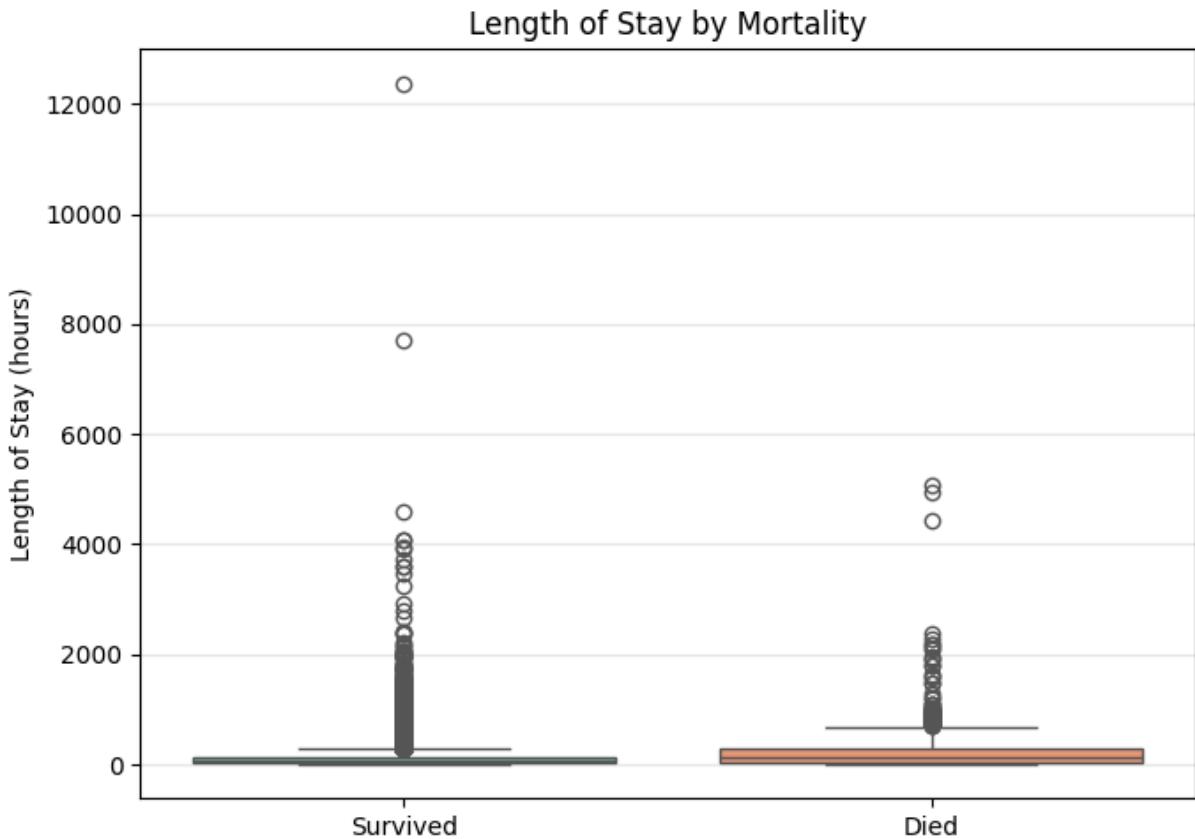
It highlights that patients who died generally had longer hospital stays, consistent with the statistical test results.

```
In [ ]: # Visualization
plt.figure(figsize=(7, 5))
sns.boxplot(x='died_in_hospital', y='los_hours', data=eda_df, palette='Set2')
plt.xticks([0, 1], ['Survived', 'Died'])
plt.title('Length of Stay by Mortality')
plt.xlabel('')
plt.ylabel('Length of Stay (hours)')
plt.grid(axis='y', alpha=0.3)
plt.tight_layout()
plt.show()
```

/tmp/ipython-input-1185102687.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='died_in_hospital', y='los_hours', data=eda_df, palette='Set2')
```



This boxplot shows that patients who **died in the hospital generally had longer lengths of stay** compared to those who survived.

The wide range and presence of extreme outliers indicate that while most patients had relatively short stays, a subset—especially among those who died—remained hospitalized much longer.

Statistical Analysis: Gender and Mortality (Chi-Squared Test)

Performing a Chi-Squared test to determine if there is a statistically significant association between patient gender and in-hospital mortality.

The Chi-Squared test was used for the analysis of Admission Type and Mortality because it's the appropriate statistical test for examining the association between two categorical variables.

Admission Type is a categorical variable (e.g., 'EW EMER.', 'ELECTIVE', 'URGENT'). In-hospital mortality is also a categorical variable (binary: Died/Survived, or 0/1). The Chi-Squared test helps us determine if the observed frequencies in each category combination (e.g., how many patients admitted via 'EW EMER.' died vs. survived) are significantly different from what would be expected if there were no association between the two variables.

Hypotheses:

- **Null Hypothesis (H0):** There is no statistically significant association between patient gender and in-hospital mortality.
- **Alternative Hypothesis (H1):** There is a statistically significant association between patient gender and in-hospital mortality.

Performing a Chi-Squared test to determine if there is a statistically significant association between patient gender and in-hospital mortality.

```
In [ ]: from scipy.stats import chi2_contingency

# Ensure mortality column is correctly identified
if 'hospital_expire_flag' in eda_df.columns:
    mortality_col = 'hospital_expire_flag'
else:
    mortality_col = 'died_in_hospital'

# Create a contingency table of Gender and Mortality
contingency_table = pd.crosstab(eda_df['gender'], eda_df[mortality_col])

print("Contingency Table (Gender vs. Mortality):")
print(contingency_table)

# Perform the Chi-Squared test
chi2, p, dof, expected = chi2_contingency(contingency_table)

print(f"\nChi-Squared Statistic: {chi2:.3f}")
print(f"P-value: {p:.4f}")
print(f"Degrees of Freedom: {dof}")

if p < 0.05:
    print("\nResult: There is a statistically significant association between gender and mortality")
else:
    print("\nResult: There is no statistically significant association between gender and mortality")
```

```
Contingency Table (Gender vs. Mortality):
hospital_expire_flag      0      1
gender
F                      25109  487
M                      23789  615
```

```
Chi-Squared Statistic: 21.811
P-value: 0.0000
Degrees of Freedom: 1
```

Result: There is a statistically significant association between gender and in-hospital mortality.

This visualization shows the mortality rate for each gender, visually supporting the findings of the Chi-Squared test.

```
In [ ]: import matplotlib.pyplot as plt
import seaborn as sns

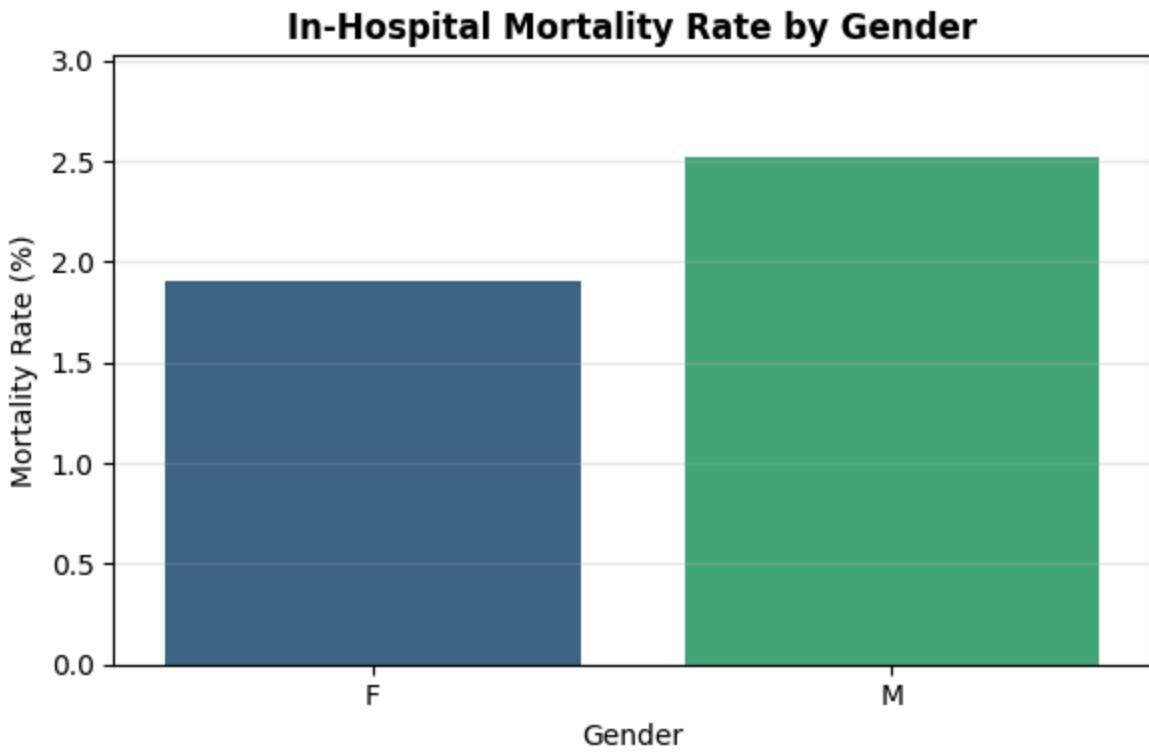
# Calculate mortality rate by gender
mortality_by_gender_plot = eda_df.groupby('gender')[mortality_col].mean().reset_index()
mortality_by_gender_plot['rate_%'] = mortality_by_gender_plot[mortality_col] * 100

plt.figure(figsize=(6, 4))
sns.barplot(x='gender', y='rate_%', data=mortality_by_gender_plot, palette='viridis')
plt.title('In-Hospital Mortality Rate by Gender', fontweight='bold')
plt.xlabel('Gender')
plt.ylabel('Mortality Rate (%)')
plt.ylim(0, mortality_by_gender_plot['rate_%'].max() * 1.2) # Add some padding
plt.grid(axis='y', alpha=0.3)
plt.tight_layout()
plt.show()
```

```
/tmp/ipython-input-2788157173.py:9: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
    sns.barplot(x='gender', y='rate_%', data=mortality_by_gender_plot, palette
='viridis')
```



Statistical Analysis: Admission Type and Mortality (Chi-Squared Test)

Performing a Chi-Squared test to determine if there is a statistically significant association between patient admission type and in-hospital mortality.

The Chi-Squared test was used for the analysis of Admission Type and Mortality because it's the appropriate statistical test for examining the association between two categorical variables.

Admission Type is a categorical variable (e.g., 'EW EMER.', 'ELECTIVE', 'URGENT'). In-hospital mortality is also a categorical variable (binary: Died/Survived, or 0/1). The Chi-Squared test helps us determine if the observed frequencies in each category combination (e.g., how many patients admitted via 'EW EMER.' died vs. survived) are significantly different from what would be expected if there were no association between the two variables.

Hypotheses:

- **Null Hypothesis (H0):** There is no statistically significant association between patient admission type and in-hospital mortality.
- **Alternative Hypothesis (H1):** There is a statistically significant association between patient admission type and in-hospital mortality.

Performing a Chi-Squared test to determine if there is a statistically significant association between patient admission type and in-hospital mortality.

```
In [ ]: from scipy.stats import chi2_contingency
import pandas as pd

# Ensure mortality column is correctly identified
if 'hospital_expire_flag' in eda_df.columns:
    mortality_col = 'hospital_expire_flag'
else:
    mortality_col = 'died_in_hospital'

# Create a contingency table of Admission Type and Mortality
contingency_table_admission = pd.crosstab(eda_df['admission_type'], eda_df[mortality_col])

print("Contingency Table (Admission Type vs. Mortality):")
print(contingency_table_admission)

# Perform the Chi-Squared test
chi2_admission, p_admission, dof_admission, expected_admission = chi2_contingency(contingency_table_admission)

print(f"\nChi-Squared Statistic: {chi2_admission:.3f}")
print(f"P-value: {p_admission:.4f}")
print(f"Degrees of Freedom: {dof_admission}")

if p_admission < 0.05:
    print("\nResult: There is a statistically significant association between admission type and mortality")
else:
    print("\nResult: There is no statistically significant association between admission type and mortality")
```

Contingency Table (Admission Type vs. Mortality):

	0	1
admission_type		
AMBULATORY OBSERVATION	613	0
DIRECT EMER.	1950	55
DIRECT OBSERVATION	2232	2
ELECTIVE	1141	6
EU OBSERVATION	11094	9
EW EMER.	15497	627
OBSERVATION ADMIT	7545	159
SURGICAL SAME DAY ADMISSION	3960	8
URGENT	4866	236

Chi-Squared Statistic: 735.674

P-value: 0.0000

Degrees of Freedom: 8

Result: There is a statistically significant association between admission type and in-hospital mortality.

```
In [ ]: import matplotlib.pyplot as plt
import seaborn as sns

# Calculate mortality rate by admission type
mortality_by_admission_plot = eda_df.groupby('admission_type')[mortality_col].sum()
mortality_by_admission_plot['rate_%'] = mortality_by_admission_plot[mortality_col] / mortality_by_admission_plot['count']

plt.figure(figsize=(10, 6))
sns.barplot(x='admission_type', y='rate_%', data=mortality_by_admission_plot)
```

```

plt.title('In-Hospital Mortality Rate by Admission Type', fontweight='bold')
plt.xlabel('Admission Type')
plt.ylabel('Mortality Rate (%)')
plt.ylim(0, mortality_by_admission_plot['rate_%'].max() * 1.2) # Add some padding at the top
plt.xticks(rotation=45, ha='right') # Rotate labels for better readability
plt.grid(axis='y', alpha=0.3)
plt.tight_layout()
plt.show()

```

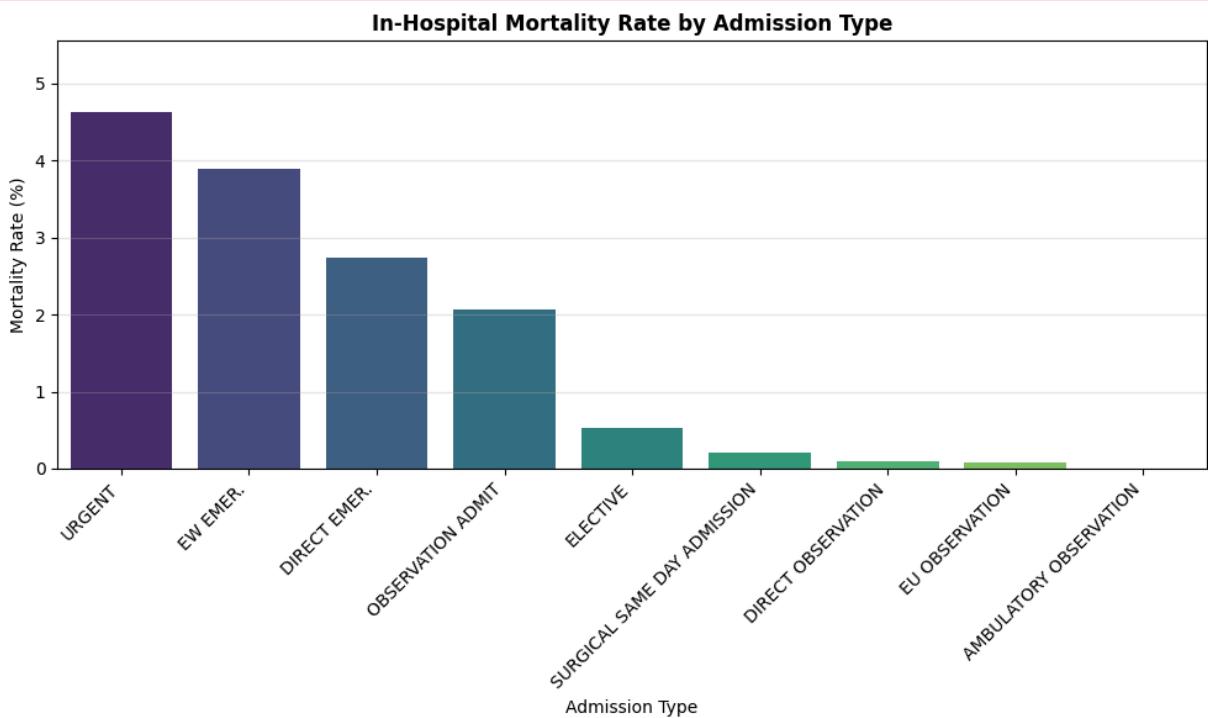
/tmp/ipython-input-495625549.py:9: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.barplot(x='admission_type', y='rate_%', data=mortality_by_admission_plot.sort_values('rate_%', ascending=False), palette='viridis')

```



Statistical Analysis: Age and Mortality (Independent t-test)

Performing an independent t-test to compare the average age of patients who died in the hospital versus those who survived.

The independent t-test is used here because we are comparing the means of a continuous variable (patient age) between two independent groups (patients who died in the hospital and patients who survived). This test helps determine if the observed difference in average age between these two groups is statistically significant or likely due to random chance.

Hypotheses:

- **Null Hypothesis (H0):** The average age of patients who died in the hospital is equal to the average age of patients who survived.
- **Alternative Hypothesis (H1):** The average age of patients who died in the hospital is not equal to the average age of patients who survived.

Performing an independent t-test to compare the average age of patients who died in the hospital versus those who survived.

```
In [ ]: from scipy.stats import ttest_ind

# Ensure mortality column is correctly identified
if 'hospital_expire_flag' in eda_df.columns:
    mortality_col = 'hospital_expire_flag'
else:
    mortality_col = 'died_in_hospital'

# Separate age data for the two groups (died vs. survived)
age_died = eda_df.loc[eda_df[mortality_col] == 1, 'anchor_age'].dropna()
age_survived = eda_df.loc[eda_df[mortality_col] == 0, 'anchor_age'].dropna()

# Perform the independent t-test (assuming unequal variances - Welch's t-test)
t_stat, p_val = ttest_ind(age_died, age_survived, equal_var=False)

print(f"T-statistic: {t_stat:.3f}")
print(f"P-value: {p_val:.4f}")

if p_val < 0.05:
    print("\nResult: There is a statistically significant difference in average age between patients who died and those who survived.")
    print(f"Average age (Died): {age_died.mean():.2f}")
    print(f"Average age (Survived): {age_survived.mean():.2f}")
else:
    print("\nResult: There is no statistically significant difference in average age between patients who died and those who survived.")

T-statistic: 30.940
P-value: 0.0000
```

Result: There is a statistically significant difference in average age between patients who died and those who survived.

Average age (Died): 70.52
 Average age (Survived): 56.67

```
In [ ]: import matplotlib.pyplot as plt
import seaborn as sns

# Ensure mortality column is correctly identified
if 'hospital_expire_flag' in eda_df.columns:
    mortality_col = 'hospital_expire_flag'
else:
    mortality_col = 'died_in_hospital'

# Visualization
plt.figure(figsize=(7, 5))
sns.boxplot(x=mortality_col, y='anchor_age', data=eda_df, palette='Set2')
plt.xticks([0, 1], ['Survived', 'Died'])
```

```

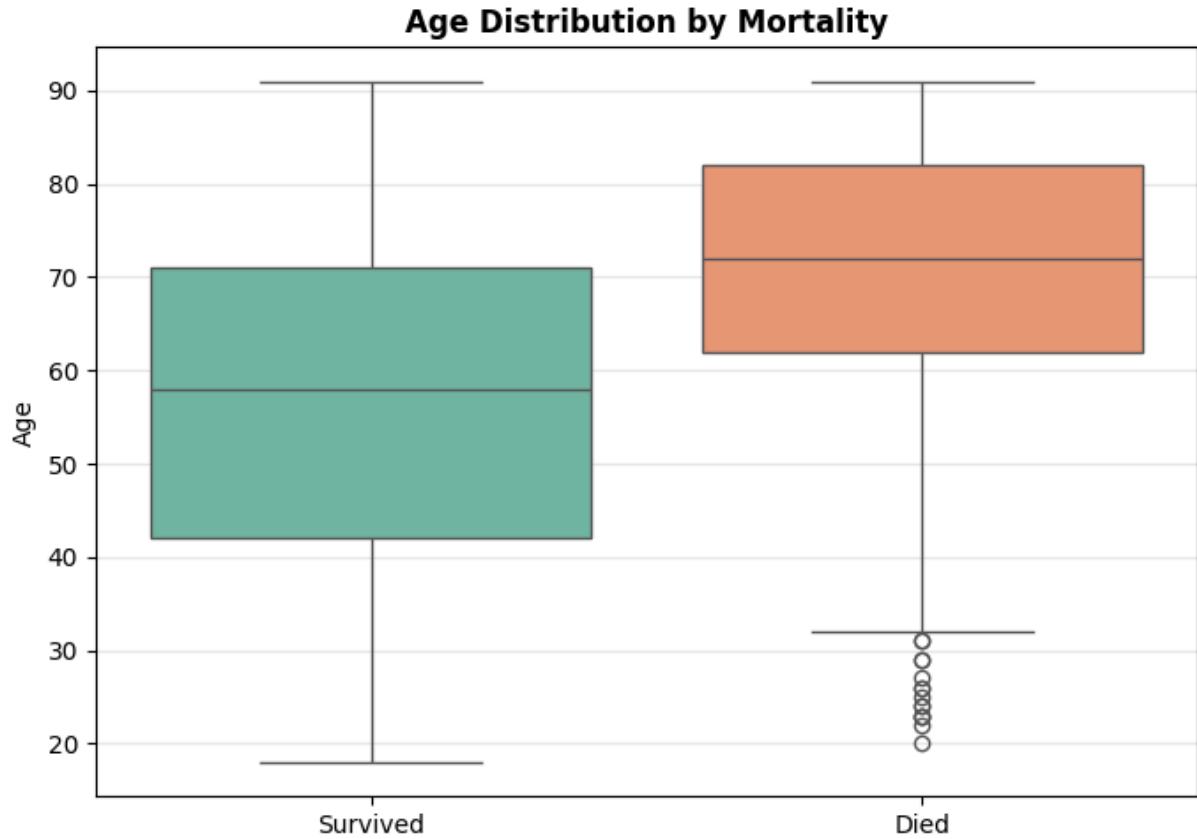
plt.title('Age Distribution by Mortality', fontweight='bold')
plt.xlabel('')
plt.ylabel('Age')
plt.grid(axis='y', alpha=0.3)
plt.tight_layout()
plt.show()

```

/tmp/ipython-input-2486593892.py:12: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x=mortality_col, y='anchor_age', data=eda_df, palette='Set2')
```



Summary of Findings and Conclusion

Based on the exploratory data analysis and the statistical tests performed, here is a summary of the key findings regarding the factors associated with in-hospital mortality:

Overall Mortality: The dataset shows an overall in-hospital mortality rate of approximately 2.20%.

Length of Stay: Patients who died in the hospital had significantly longer lengths of stay compared to those who survived (as demonstrated by the Welch's t-test with a highly significant p-value). The correlation analysis also showed a **positive correlation** between `los_hours` and `died_in_hospital`, supporting this finding.

Age: There is a statistically significant difference in average age between patients who died and those who survived. Patients who died were, on average, older (mean age ~70.5) than those who survived (mean age ~56.7) (as shown by the independent t-test with a highly significant p-value). The mortality rate increases with age, with the highest rate observed in the 75+ age group. The correlation analysis showed a **positive correlation** between `anchor_age` and `died_in_hospital`.

Gender: The Chi-Squared test revealed a statistically significant association between patient gender and in-hospital mortality. While the difference in mortality rates between genders is not as pronounced as with age or length of stay, the analysis suggests that gender is a statistically relevant factor. (Visualizations show a slightly higher mortality rate for males compared to females).

Admission Type: The Chi-Squared test demonstrated a highly statistically significant association between admission type and in-hospital mortality. Certain admission types, particularly `URGENT` and `EW EMER.`, are associated with substantially higher mortality rates compared to elective or observation-based admissions. This highlights the acuity and severity of illness upon arrival as a critical factor.

Race/Ethnicity and Insurance: While the Chi-Squared tests for these variables were not explicitly performed in the provided cells, the mortality rate analysis by group showed variations, notably a higher mortality rate in the `UNKNOWN` race category and among patients with `OTHER` and `MEDICARE` insurance types. Further statistical testing would be needed to confirm the significance of these associations.

Text Data Completeness: A significant portion of records have missing discharge and radiology text data (22.3% and 17.5% respectively). While this doesn't directly impact the current mortality analysis based on structured data, it is a crucial data quality issue for the downstream RAG model, which relies heavily on this unstructured text.

Conclusion:

The exploratory data analysis and statistical tests strongly indicate that patient age, length of hospital stay, gender, and admission type are significant factors associated with in-hospital mortality in this dataset. Older patients, those with longer hospitalizations, and those admitted through emergency pathways (`EW EMER.`, `URGENT`) have a higher risk of mortality. While gender shows a statistically significant association, the practical difference in mortality rate is less pronounced compared to age and admission type. The availability of unstructured text data is a notable limitation for future RAG model development. These findings are crucial for building a predictive model and for interpreting the outputs of the Medical Evidence Synthesizer.